

Prelink

Jakub Jelínek
Red Hat, Inc.
jakub@redhat.com

December 10, 2003

Abstract

Prelink is a tool designed to speed up dynamic linking of ELF programs on various Linux architectures. It speeds up start up of OpenOffice.org 1.1 by 1.8s from 5.5s on 651MHz Pentium III.

1 Preface

1 In 1995, Linux changed its binary format from `a.out` to `ELF`. The `a.out` binary format was very inflexible and shared
2 libraries were pretty hard to build. Linux's shared libraries in `a.out` are position dependent and each had to be given a
3 unique virtual address space slot at link time. Maintaining these assignments was pretty hard even when there were just
4 a few shared libraries, there used to be a central address registry maintained by humans in form of a text file, but it is
5 certainly impossible to do these days when there are thousands of different shared libraries and their size, version and
6 exported symbols are constantly changing. On the other side, there was just minimum amount of work the dynamic
7 linker had to do in order to load these shared libraries, as relocation handling and symbol lookup was only done at link
8 time. The dynamic linker used the `useLib` system call which just mapped the named library into the address space
9 (with no segment or section protection differences, the whole mapping was writable and executable).

10 The `ELF`¹ binary format is one of the most flexible binary formats, its shared libraries are easy to build and there is no
11 need for a central assignment of virtual address space slots. Shared libraries are position independent and relocation
12 handling and symbol lookup are done partly at the time the executable is created and partly at runtime. Symbols in
13 shared libraries can be overridden at runtime by preloading a new shared library defining those symbols or without
14 relinking an executable by adding symbols to a shared library which is searched up earlier during symbol lookup or
15 by adding new dependent shared libraries to a library used by the program. All these improvements have their price,
16 which is a slower program startup, more non-shareable memory per process and runtime cost associated with position
17 independent code in shared libraries.

18 Program startup of `ELF` programs is slower than startup of `a.out` programs with shared libraries, because the dynamic
19 linker has much more work to do before calling program's entry point. The cost of loading libraries is just slightly
20 bigger, as `ELF` shared libraries have typically separate read-only and writable segments, so the dynamic linker has to
21 use different memory protection for each segment. The main difference is in relocation handling and associated symbol
22 lookup. In the `a.out` format there was no relocation handling or symbol lookup at runtime. In `ELF`, this cost is much
23 more important today than it used to be during `a.out` to `ELF` transition in Linux, as especially GUI programs keep
24 constantly growing and start to use more and more shared libraries. 5 years ago programs using more than 10 shared
25 libraries were very rare, these days most of the GUI programs link against around 40 or more shared and in extreme
26 cases programs use even more than 90 shared libraries. Every shared library adds its set of dynamic relocations to
27 the cost and enlarges symbol search scope, so in addition to doing more symbol lookups, each symbol lookup the
28 application has to perform is on average more expensive. Another factor increasing the cost is the length of symbol
29 names which have to be compared when finding symbol in the symbol hash table of a shared library. C++ libraries
30 tend to have extremely long symbol names and unfortunately the new C++ ABI puts namespaces and class names first
31 and method names last in the mangled names, so often symbol names differ only in last few bytes of very long names.

32 Every time a relocation is applied the entire memory page containing the address which is written to must be loaded
33 into memory. The operating system does a copy-on-write operation which also has the consequence that the physical

¹As described in generic ABI document [1] and various processor specific ABI supplements [2], [3], [4], [5], [6], [7], [8].

³⁴ memory of the memory page cannot anymore be shared with other processes. With ELF, typically all of program's
³⁵ Global Offset Table, constants and variables containing pointers to objects in shared libraries, etc. are written into
³⁶ before the dynamic linker passes control over to the program.

³⁷ On most architectures (with some exceptions like AMD64 architecture) position independent code requires that one
³⁸ register needs to be dedicated as PIC register and thus cannot be used in the functions for other purposes. This
³⁹ especially degrades performance on register-starved architectures like IA-32. Also, there needs to be some code to
⁴⁰ set up the PIC register, either invoked as part of function prologues, or when using function descriptors in the calling
⁴¹ sequence.

⁴² Prelink is a tool which (together with corresponding dynamic linker and linker changes) attempts to bring back some
⁴³ of the a.out advantages (such as the speed and less COW'd pages) to the ELF binary format while retaining all of
⁴⁴ its flexibility. In a limited way it also attempts to decrease number of non-shareable pages created by relocations.
⁴⁵ Prelink works closely with the dynamic linker in the GNU C library, but probably it wouldn't be too hard to port it
⁴⁶ to some other ELF using platforms where the dynamic linker can be modified in similar ways.

2 Caching of symbol lookup results

⁴⁷ Program startup can be speeded up by caching of symbol lookup results². Many shared libraries need more than one
⁴⁸ lookup of a particular symbol. This is especially true for C++ shared libraries, where e.g. the same method is present in
⁴⁹ multiple virtual tables or RTTI data structures. Traditionally, each ELF section which needs dynamic relocations has an
⁵⁰ associated .rela* or .rel* section (depending on whether the architecture is defined to use RELA or REL relocations).

⁵¹ The relocations in those sections are typically sorted by ascending r_offset values. Symbol lookups are usually the
⁵² most expensive operation during program startup, so caching the symbol lookups has potential to decrease time spent
⁵³ in the dynamic linker. One way to decrease the cost of symbol lookups is to create a table with the size equal to number
⁵⁴ of entries in dynamic symbol table (.dynsym) in the dynamic linker when resolving a particular shared library, but that
⁵⁵ would in some cases need a lot of memory and some time spent in initializing the table. Another option would be to
⁵⁶ use a hash table with chained lists, but that needs both extra memory and would also take extra time for computation of
⁵⁷ the hash value and walking up the chains when doing new lookups. Fortunately, neither of this is really necessary if we
⁵⁸ modify the linker to sort relocations so that relocations against the same symbol are adjacent. This has been done first
⁵⁹ in the Sun linker and dynamic linker, so the GNU linker and dynamic linker use the same ELF extensions and linker
⁶⁰ flags. Particularly, the following new ELF dynamic tags have been introduced:

```
61 #define DT_RELACOUNT 0x6fffffff9
62 #define DT_RELCOUNT 0x6fffffff9a
```

⁶³ New options -z combreloc and -z nocombrelod have been added to the linker. The latter causes the previous
⁶⁴ linker behavior, i.e. each section requiring relocations has a corresponding relocation section, which is sorted by
⁶⁵ ascending r_offset. -z combreloc³ instructs the linker to create just one relocation section for dynamic relocations
⁶⁶ other than symbol jump table (PLT) relocations. This single relocation section (either .rela.dyn or .rel.dyn) is
⁶⁷ sorted, so that relative relocations come first (sorted by ascending r_offset), followed by other relocations, sorted
⁶⁸ again by ascending r_offset. If more relocations are against the same symbol, they immediately follow the first
⁶⁹ relocation against that symbol with lowest r_offset.⁴ The number of relative relocations at the beginning of the
⁷⁰ section is stored in the DT_RELACOUNT resp. DT_RELCOUNT dynamic tag.

⁷¹ The dynamic linker can use the new dynamic tag for two purposes. If the shared library is successfully mapped at the
⁷² same address as the first PT_LOAD segment's virtual address, the load offset is zero and the dynamic linker can avoid all
⁷³ the relative relocations which would just add zero to various memory locations. Normally shared libraries are linked
⁷⁴ with first PT_LOAD segment's virtual address set to zero, so the load offset is non-zero. This can be changed through
⁷⁵ a linker script or by using a special prelink option --reloc-only to change the base address of a shared library.
⁷⁶ All prelinked shared libraries have non-zero base address as well. If the load offset is non-zero, the dynamic linker
⁷⁷ can still make use of this dynamic tag, as relative relocation handling is typically way simpler than handling other

²Initially, this has been implemented in the prelink tool and glibc dynamic linker, where prelink was sorting relocation sections of existing executables and shared libraries. When this has been implemented in the linker as well and most executables and shared libraries are already built with -z combreloc, the code from prelink has been removed, as it was no longer needed for most objects and just increasing the tool's complexity.

³-z combreloc is the default in GNU linker versions 2.13 and later.

⁴In fact the sorting needs to take into account also the type of lookup. Most of the relocations will resolve to a PLT slot in the executable if there is one for the lookup symbol, because the executable might have a pointer against that symbol without any dynamic relocations. But e.g. relocations used for the PLT slots must avoid these.

78 relocations (since symbol lookup is not necessary) and thus it can handle all relative relocations in a tight loop in one
79 place and then handle the remaining relocations with the fully featured relocation handling routine. The second and
80 more important point is that if relocations against the same symbol are adjacent, the dynamic linker can use a cache
81 with single entry.

82 The dynamic linker in `glibc`, if it sees `statistics` as part of the `LD_DEBUG` environment variable, displays statistics
83 which can show how useful this optimization is. Let's look at some big C++ application, e.g. `konqueror`. If not using
84 the cache, the statistics looks like this:

```
85 18000: runtime linker statistics:  
86 18000:     total startup time in dynamic loader: 270886059 clock cycles  
87 18000:             time needed for relocation: 266364927 clock cycles (98.3%)  
88 18000:             number of relocations: 79067  
89 18000:             number of relocations from cache: 0  
90 18000:             number of relative relocations: 31169  
91 18000:             time needed to load objects: 4203631 clock cycles (1.5%)
```

92 This program run is with hot caches, on non-prelinked system, with lazy binding. The numbers show that the dynamic
93 linker spent most of its time in relocation handling and especially symbol lookups. If using symbol lookup cache, the
94 numbers look different:

```
95 18013:     total startup time in dynamic loader: 132922001 clock cycles  
96 18013:             time needed for relocation: 128399659 clock cycles (96.5%)  
97 18013:             number of relocations: 25473  
98 18013:             number of relocations from cache: 53594  
99 18013:             number of relative relocations: 31169  
100 18013:             time needed to load objects: 4202394 clock cycles (3.1%)
```

101 On average, for one real symbol lookup there were two cache hits and total time spent in the dynamic linker decreased
102 by 50%.

3 Prelink design

103 Prelink was designed, so that it requires as few ELF extensions as possible. It should not be tied to a particular
104 architecture, but should work on all ELF architectures. During program startup it should avoid all symbol lookups
105 which, as has been shown above, are very expensive. It needs to work in an environment where shared libraries and
106 executables are changing from time to time, whether it is because of security updates or feature enhancements. It
107 should avoid big code duplication between the dynamic linker and the tool. And prelinked shared libraries need to be
108 usable even in non-prelinked executables, or when one of the shared libraries is upgraded and the prelinking of the
109 executable has not been updated.

110 To minimize the number of performed relocations during startup, the shared libraries (and executables) need to be
111 relocated already as much as possible. For relative relocations this means the library needs to be loaded always at
112 the same base address, for other relocations this means all shared libraries with definitions those relocations resolve
113 to (often this includes all shared libraries the library or executable depends on) must always be loaded at the same
114 addresses. ELF executables (with the exception of *Position Independent Executables*) have their load address fixed
115 already during linking. For shared libraries, prelink needs something similar to a .out registry of virtual address
116 space slots. Maintaining such registry across all installations wouldn't scale well, so prelink instead assigns these
117 virtual address space slots on the fly after looking at all executables it is supposed to speed up and all their dependent
118 shared libraries. The next step is to actually relocate shared libraries to the assigned base address.

119 When this is done, the actual prelinking of shared libraries can be done. First, all dependent shared libraries need to be
120 prelinked (prelink doesn't support circular dependencies between shared libraries, will just warn about them instead
121 of prelinking the libraries in the cycle), then for each relocation in the shared library prelink needs to look up the
122 symbol in natural symbol search scope of the shared library (the shared library itself first, then breadth first search of
123 all dependent shared libraries) and apply the relocation to the symbol's target section. The symbol lookup code in the

¹²⁴ dynamic linker is quite complex and big, so to avoid duplicating all this, `prelink` has chosen to use dynamic linker to
¹²⁵ do the symbol lookups. Dynamic linker is told via a special environment variable it should print all performed symbol
¹²⁶ lookups and their type and `prelink` reads this output through a pipe. As one of the requirements was that prelinked
¹²⁷ shared libraries must be usable even for non-prelinked executables (duplicating all shared libraries so that there are
¹²⁸ pristine and prelinked copies would be very unfriendly to RAM usage), `prelink` has to ensure that by applying the
¹²⁹ relocation no information is lost and thus relocation processing can be cheaply done at startup time of non-prelinked
¹³⁰ executables. For RELA architectures this is easier, because the content of the relocation's target memory is not needed
¹³¹ when processing the relocation.⁵ For REL architectures this is not the case. `prelink` attempts some tricks described
¹³² later and if they fail, needs to convert the REL relocation section to RELA format where addend is stored in the relocation
¹³³ section instead of relocation target's memory.

¹³⁴ When all shared libraries an executable (directly or indirectly) depends on are prelinked, relocations in the executable
¹³⁵ are handled similarly to relocations in shared libraries. Unfortunately, not all symbols resolve the same when looked up
¹³⁶ in a shared library's natural symbol search scope (i.e. as it is done at the time the shared library is prelinked) and when
¹³⁷ looked up in application's global symbol search scope. Such symbols are herein called *conflicts* and the relocations
¹³⁸ against those symbols *conflicting relocations*. Conflicts depend on the executable, all its shared libraries and their
¹³⁹ respective order. They are only computable for the shared libraries linked to the executable (libraries mentioned in
¹⁴⁰ DT_NEEDED dynamic tags and shared libraries they transitively need). The set of shared libraries loaded via `dlopen(3)`
¹⁴¹ cannot be predicted by `prelink`, neither can the order in which this happened, nor the time when they are unloaded.
¹⁴² When the dynamic linker prints symbol lookups done in the executable, it also prints conflicts. `Prelink` then takes all
¹⁴³ relocations against those symbols and builds a special RELA section with conflict fixups and stores it into the prelinked
¹⁴⁴ executable. Also a list of all dependent shared libraries in the order they appear in the symbol search scope, together
¹⁴⁵ with their checksums and times of prelinking is stored in another special section.

¹⁴⁶ The dynamic linker first checks if it is itself prelinked. If yes, it can avoid its preliminary relocation processing (this
¹⁴⁷ one is done with just the dynamic linker itself in the search scope, so that all routines in the dynamic linker can be
¹⁴⁸ used easily without too many limitations). When it is about to start a program, it first looks at the library list section
¹⁴⁹ created by `prelink` (if any) and checks whether they are present in symbol search scope in the same order, none
¹⁵⁰ have been modified since prelinking and that there aren't any new shared libraries loaded either. If all these conditions
¹⁵¹ are satisfied, prelinking can be used. In that case the dynamic linker processes the fixup section and skips all normal
¹⁵² relocation handling. If one or more of the conditions are not met, the dynamic linker continues with normal relocation
¹⁵³ processing in the executable and all shared libraries.

4 Collecting executables and libraries which should be prelinked

¹⁵⁴ Before the actual work can start the `prelink` tool needs to collect the filenames of executables and libraries it is sup-
¹⁵⁵ posed to prelink. It doesn't make any sense to prelink a shared library if no executable is linked against it because
¹⁵⁶ the prelinking information will not be used anyway. Furthermore, when `prelink` needs to do a REL TO RELA con-
¹⁵⁷ version of relocation sections in the shared library (see later) or when it needs to convert SHT_NOBITS PLT section to
¹⁵⁸ SHT_PROGBITS, a prelinked shared library might grow in size and so prelinking is only desirable if it will speed up
¹⁵⁹ startup of some program. The only change which might be useful even for shared libraries which are never linked
¹⁶⁰ against, only loaded using `dlopen`, is relocating to a unique address. This is useful if there are many relative relo-
¹⁶¹ locations and there are pages in the shared library's writable segment which are never written into with the exception
¹⁶² of those relative relocations. Such shared libraries are rare, so `prelink` doesn't handle these automatically, instead
¹⁶³ the administrator or developer can use `prelink --reloc-only=ADDRESS` to relocate it manually. Prelinking an
¹⁶⁴ executable requires all shared libraries it is linked against to be prelinked already.

¹⁶⁵ `Prelink` has two main modes in which it collects filenames. One is *incremental prelinking*, where `prelink` is
¹⁶⁶ invoked without the `-a` option. In this mode, `prelink` queues for prelinking all executables and shared libraries given
¹⁶⁷ on the command line, all executables in directory trees specified on the command line, and all shared libraries those
¹⁶⁸ executables and shared libraries are linked against. For the reasons mentioned earlier a shared library is queued only if
¹⁶⁹ a program is linked with it or the user tells the tool to do it anyway by explicitly mentioning it on the command line.
¹⁷⁰ The second mode is *full prelinking*, where the `-a` option is given on the command line. This in addition to incremental
¹⁷¹ prelinking queues all executables found in directory trees specified in `prelink.conf` (which typically includes all or
¹⁷² most directories where system executables are found). For each directory subtree in the config file the user can specify
¹⁷³ whether symbolic links to places outside of the tree are to be followed or not and whether searching should continue
¹⁷⁴ even across filesystem boundaries.

⁵Relative relocations on certain RELA architectures use relocation target's memory, either alone or together with `r_addend` field.

175 There is also an option to blacklist some executables or directory trees so that the executables or anything in the
176 directory trees will not be prelinked. This can be specified either on the command line or in the config file.

177 Prelink will not attempt to change executables which use a non-standard dynamic linker ⁶ for security reasons,
178 because it actually needs to execute the dynamic linker for symbol lookup and it needs to avoid executing some random
179 unknown executable with the permissions with which `prelink` is run (typically `root`, with the permissions at least
180 for changing all executables and shared libraries in the system). The administrator should ensure that `prelink.conf`
181 doesn't contain world-writable directories and such directories are not given to the tool on the command line either, but
182 the tool should be distrustful of the objects nevertheless.

183 Also, `prelink` will not change shared libraries which are not specified directly on the command line or located in the
184 directory trees specified on the command line or in the config file. This is so that e.g. `prelink` doesn't try to change
185 shared libraries on shared networked filesystems, or at least it is possible to configure the tool so that it doesn't do it.

186 For each executable and shared library it collects, `prelink` executes the dynamic linker to list all shared libraries it
187 depends on, checks if it is already prelinked and whether any of its dependencies changed. Objects which are already
188 prelinked and have no dependencies which changed don't have to be prelinked again (with the exception when e.g.
189 virtual address space layout code finds out it needs to assign new virtual address space slots for the shared library or
190 one of its dependencies). Running the dynamic linker to get the symbol lookup information is a quite costly operation
191 especially on systems with many executables and shared libraries installed, so `prelink` offers a faster `-q` mode. In
192 all modes, `prelink` stores modification and change times of each shared library and executable together with all
193 object dependencies and other information into `prelink.cache` file. When prelinking in `-q` mode, it just compares
194 modification and change times of the executables and shared libraries (and all their dependencies). Change time is
195 needed because `prelink` preserves modification time when prelinking (as well as permissions, owner and group). If
196 the times match, it assumes the file has not changed since last prelinking. Therefore the file can be skipped if it is
197 already prelinked and none of the dependencies changed. If any time changed or one of the dependencies changed, it
198 invokes the dynamic linker the same way as in normal mode to find out real dependencies, whether it has been prelinked
199 or not etc. The collecting phase in normal mode can take a few minutes, while in quick mode usually takes just a few
200 seconds, as the only operation it does is it calls just lots of `stat` system calls.

5 Assigning virtual address space slots

201 Prelink has to ensure at least that for all successfully prelinked executables all shared libraries they are (transitively)
202 linked against have non-overlapping virtual address space slots (furthermore they cannot overlap with the virtual ad-
203 dress space range used by the executable itself, its `brk` area, typical stack location and `ld.so.cache` and other files
204 mmaped by the dynamic linker in early stages of dynamic linking (before all dependencies are mmaped)). If there were
205 any overlaps, the dynamic linker (which mmapps the shared libraries at the desired location without `MAP_FIXED` mmap
206 flag so that it is only soft requirement) would not manage to mmap them at the assigned locations and the prelinking
207 information would be invalidated (the dynamic linker would have to do all normal relocation handling and symbol
208 lookups). Executables are linked against very wide variety of shared library combinations and that has to be taken into
209 account.

210 The simplest approach is to sort shared libraries by descending usage count (so that most often used shared libraries
211 like the dynamic linker, `libc.so` etc. are close to each other) and assign them consecutive slots starting at some
212 architecture specific base address (with a page or two in between the shared libraries to allow for a limited growth of
213 shared libraries without having to reposition them). Prelink has to find out which shared libraries will need a `REL` to
214 `RELA` conversion of relocation sections and for those which will need the conversion count with the increased size of
215 the library's loadable segments. This is `prelink` behavior without `-m` and `-R` options.

216 The architecture specific base address is best located a few megabytes above the location where `mmap` with `NULL` first
217 argument and without `MAP_FIXED` starts allocating memory areas (in Linux this is the value of `TASK_UNMAPPED_BASE`
218 macro). ⁷ The reason for not starting to assign addresses in `prelink` immediately at `TASK_UNMAPPED_BASE` is that
219 `ld.so.cache` and other mappings by the dynamic linker will end up in the same range and could overlap with the
220 shared libraries. Also, if some application uses `dlopen` to load a shared library which has been prelinked, ⁸ those

⁶Standard dynamic linker path is hardcoded in the executable for each architecture. It can be overridden from the command line, but only with one dynamic linker name (normally, multiple standard dynamic linkers are used when prelinking mixed architecture systems).

⁷`TASK_UNMAPPED_BASE` has been chosen on each platform so that there is enough virtual memory for both the `brk` area (between executable's end and this memory address) and `mmap` area (between this address and bottom of stack).

⁸Typically this is because some other executable is linked against that shared library directly.

221 few megabytes above `TASK_UNMAPPED_BASE` increase the probability that the stack slot will be still unused (it can
222 clash with e.g. non-prelinked shared libraries loaded by `dlopen` earlier⁹ or other kinds of `mmap` calls with `NULL` first
223 argument like `malloc` allocating big chunks of memory, mmaping of locale database, etc.).

224 This simplest approach is unfortunately problematic on 32-bit (or 31-bit) architectures where the total virtual address
225 space for a process is somewhere between 2GB (S/390) and almost 4GB (Linux IA-32 4GB/4GB kernel split, AMD64
226 running 32-bit processes, etc.). Typical installations these days contain thousands of shared libraries and if each of
227 them is given a unique address space slot, on average executables will have pretty sparse mapping of its shared libraries
228 and there will be less contiguous virtual memory for application's own use¹⁰.

229 Prelink has a special mode, turned on with `-m` option, in which it computes what shared libraries are ever loaded
230 together in some executable (not considering `dlopen`). If two shared libraries are ever loaded together, `prelink`
231 assigns them different virtual address space slots, but if they never appear together, it can give them overlapping
232 addresses. For example applications using `KDE` toolkit link typically against many `KDE` shared libraries, programs
233 written using the `Gtk+` toolkit link typically against many `Gtk+` shared libraries, but there are just very few programs
234 which link against both `KDE` and `Gtk+` shared libraries, and even if they do, they link against very small subset of
235 those shared libraries. So all `KDE` shared libraries not in that subset can use overlapping addresses with all `Gtk+` shared
236 libraries but the few exceptions. This leads to considerably smaller virtual address space range used by all prelinked
237 shared libraries, but it has its own disadvantages too. It doesn't work too well with incremental prelinking, because then
238 not all executables are investigated, just those which are given on `prelink`'s command line. `Prelink` also considers
239 executables in `prelink.cache`, but it has no information about executables which have not been prelinked yet. If
240 a new executable, which links against some shared libraries which never appeared together before, is prelinked later,
241 `prelink` has to assign them new, non-overlapping addresses. This means that any executables, which linked against
242 the library that has been moved and re-prelinked, need to be prelinked again. If this happened during incremental
243 prelinking, `prelink` will fix up only the executables given on the command line, leaving other executables untouched.
244 The untouched executables would not be able to benefit from prelinking anymore.

245 Although with the above two layout schemes shared library addresses can vary slightly between different hosts running
246 the same distribution (depending on the exact set of installed executables and libraries), especially the most often used
247 shared libraries will have identical base addresses on different computers. This is often not desirable for security
248 reasons, because it makes it slightly easier for various exploits to jump to routines they want. Standard Linux kernels
249 assign always the same addresses to shared libraries loaded by the application at each run, so with these kernels
250 `prelink` doesn't make things worse. But there are kernel patches, such as Red Hat's `Exec-Shield`, which randomize
251 memory mappings on each run. If shared libraries are prelinked, they cannot be assigned different addresses on each
252 run (prelinking information can be only used to speed up startup if they are mapped at the base addresses which was
253 used during prelinking), which means prelinking might not be desirable on some edge servers. `Prelink` can assign
254 different addresses on different hosts though, which is almost the same as assigning random addresses on each run for
255 long running processes such as daemons. Furthermore, the administrator can force full prelinking and assignment of
256 new random addresses every few days (if he is also willing to restart the services, so that the old shared libraries and
257 executables don't have to be kept in memory).

258 To assign random addresses `prelink` has the `-R` option. This causes a random starting address somewhere in the
259 architecture specific range in which shared libraries are assigned, and minor random reshuffling in the queue of shared
260 libraries which need address assignment (normally it is sorted by descending usage count, with randomization shared
261 libraries which are not very far away from each other in the sorted list can be swapped). The `-R` option should work
262 orthogonally to the `-m` option.

263 Some architectures have special further requirements on shared library address assignment. On 32-bit PowerPC, if
264 shared libraries are located close to the executable, so that everything fits into 32MB area, `PLT` slots resolving to those
265 shared libraries can use the branch relative instruction instead of more expensive sequences involving memory load and
266 indirect branch. If shared libraries are located in the first 32MB of address space, `PLT` slots resolving to those shared
267 libraries can use the branch absolute instruction (but already `PLT` slots in those shared libraries resolving to addresses in
268 the executable cannot be done cheaply). This means for optimization `prelink` should assign addresses from a 24MB
269 region below the executable first, assuming most of the executables are smaller than those remaining 8MB. `prelink`
270 assigns these from higher to lower addresses. When this region is full, `prelink` starts from address 0x40000¹¹ up

⁹If shared libraries have first `PT_LOAD` segment's virtual address zero, the kernel typically picks first empty slot above `TASK_UNMAPPED_BASE` big enough for the mapping.

¹⁰Especially databases look these days for every byte of virtual address space on 32-bit architectures.

¹¹To leave some pages unmapped to catch NULL pointer dereferences.

271 till the bottom of the first area. Only when all these areas are full, `prelink` starts picking addresses high above the
272 executable, so that sufficient space is left in between to leave room for `brk`. When `-R` option is specified, `prelink`
273 needs to honor it, but in a way which doesn't totally kill this optimization. So it picks up a random start base within
274 each of the 3 regions separately, splitting them into 6 regions.

275 Another architecture which needs to be handled specially is IA-32 when using `Exec-Shield`. The IA-32 architecture
276 doesn't have an bit to disable execution for each page, only for each segment. All readable pages are normally exe-
277 cutable. This means the stack is usually executable, as is memory allocated by `malloc`. This is undesirable for security
278 reasons, exploits can then overflow a buffer on the stack to transfer control to code it creates on the stack. Only very
279 few programs actually need an executable stack. For example programs using GCC trampolines for nested functions
280 need it or when an application itself creates executable code on the stack and calls it. `Exec-Shield` works around this
281 IA-32 architecture deficiency by using a separate code segment, which starts at address 0 and spans address space until
282 its limit, highest page which needs to be executable. This is dynamically changed when some page with higher address
283 than the limit needs to be executable (either because of `mmap` with `PROT_EXEC` bit set, or `mprotect` with `PROT_EXEC`
284 of an existing mapping). This kind of protection is of course only effective if the limit is as low as possible. The
285 kernel tries to put all new mappings with `PROT_EXEC` set and NULL address low. If possible into *ASCII Shield* area
286 (first 16MB of address space), if not, at least below the executable. If `prelink` detects `Exec-Shield`, it tries to do
287 the same as kernel when assigning addresses, i.e. prefers to assign addresses in *ASCII Shield* area and continues with
288 other addresses below the program. It needs to leave first 1MB plus 4KB of address space unallocated though, because
289 that range is often used by programs using `vm86` system call.

6 Relocation of libraries

290 When a shared library has a base address assigned, it needs to be relocated so that the base address is equal to the first
291 `PT_LOAD` segment's `p_vaddr`. The effect of this operation should be bitwise identical as if the library were linked with
292 that base address originally. That is, the following scripts should produce identical output:

```
293 $ gcc -g -shared -o libfoo.so.1.0.0 -Wl,-h,libfoo.so.1 \
294     input1.o input2.o somelib.a
295 $ prelink --reloc-only=0x54321000 libfoo.so.1.0.0
```

Listing 0: Script to relocate a shared library after linking using `prelink`

296 and:

```
297 $ gcc -shared -Wl,--verbose 2>&1 > /dev/null \
298 | sed -e '/^=====,/^\=====/!d' \
299   -e '/^====/d;s/0\(+ SIZEOF_HEADERS\)/0x54321000\1/' \
300   > libfoo.so.lds
301 $ gcc -Wl,-T,libfoo.so.lds -g -shared -o libfoo.so.1.0.0 \
302   -Wl,-h,libfoo.so.1 input1.o input2.o somelib.a
```

Listing 1: Script to link a shared library at non-standard base

303 The first script creates a normal shared library with the default base address 0 and then uses `prelink`'s special mode
304 when it just relocates a library to a given address. The second script first modifies a built-in GNU linker script for
305 linking of shared libraries, so that the base address is the one given instead of zero and stores it into a temporary file.
306 Then it creates a shared library using that linker script.

307 The relocation operation involves mostly adding the difference between old and new base address to all `ELF` fields
308 which contain values representing virtual addresses of the shared library (or in the program header table also represent-
309 ing physical addresses). File offsets need to be unmodified. Most places where the adjustments need to be done are
310 clear, `prelink` just has to watch `ELF` spec to see which fields contain virtual addresses.

311 One problem is with absolute symbols. `Prelink` has no way to find out if an absolute symbol in a shared library is
312 really meant as absolute and thus not changing during relocation, or if it is an address of some place in the shared

313 library outside of any section or on their edge. For instance symbols created in the GNU linker's script outside of
314 section directives have all SHN_ABS section, yet they can be located in the library (e.g. symbolfoo = .) or they can
315 be absolute (e.g. symbolbar = 0x12345000). This distinction is lost at link time. But the dynamic linker when
316 looking up symbols doesn't make any distinction between them, all addresses during dynamic lookup have the load
317 offset added to it. Prelink chooses to relocate any absolute symbols with value bigger than zero, that way prelink
318 --reloc-only gets bitwise identical output with linking directly at the different base in almost all real-world cases.
319 Thread Local Storage symbols (those with STT_TLS type) are never relocated, as their values are relative to start of
320 shared library's thread local area.

321 When relocating the dynamic section there are no bits which tell if a particular dynamic tag uses d.un.d_ptr (which
322 needs to be adjusted) or d.un.d_val (which needs to be left as is). So prelink has to hardcode a list of well known
323 architecture independent dynamic tags which need adjusting and have a hook for architecture specific dynamic tag
324 adjustment. Sun came up with DTADDRNGLO to DTADDRNGHI and DTVALRNGLO to DTVALRNGHI dynamic tag
325 number ranges, so at least as long as these ranges are used for new dynamic tags prelink can relocate correctly even
326 without listing them all explicitly.

327 When relocating .rela.* or .rel.* sections, which is done in architecture specific code, relative relocations and
328 on .got.plt using architectures also PLT relocations typically need an adjustment. The adjustment needs to be done
329 in either r.addend field of the ElfNN_Rela structure, in the memory pointed by r.offset, or in both locations. On
330 some architectures what needs adjusting is not even the same for all relative relocations. Relative relocations against
331 some sections need to have r.addend adjusted while others need to have memory adjusted. On many architectures,
332 first few words in GOT are special and some of them need adjustment.

333 The hardest part of the adjustment is handling the debugging sections. These are non-allocated sections which typically
334 have no corresponding relocation section associated with them. Prelink has to match the various debuggers in what
335 fields it adjusts and what are skipped. As of this writing prelink should handle DWARF 2 [15] standard as corrected
336 (and extended) by DWARF 3 draft [16], Stabs [17] with GCC extensions and Alpha or MIPS Mdebug.

337 DWARF 2 debugging information involves many separate sections, each of them with a unique format which needs
338 to be relocated differently. For relocation of the .debug_info section compilation units prelink has to parse the
339 corresponding part of the .debug_abbrev section, adjust all values of attributes that are using the DW_FORM_addr
340 form and adjust embedded location lists. .debug_ranges and .debug_loc section portions depend on the exact place
341 in .debug_info section from which they are referenced, so that prelink can keep track of their base address. DWARF
342 debugging format is very extendable, so prelink needs to be very conservative when it sees unknown extensions.
343 It needs to fail prelinking instead of silently break debugging information if it sees an unknown .debug_* section,
344 unknown attribute form or unknown attribute with one of the DW_FORM_block* forms, as they can potentially embed
345 addresses which would need adjustment.

346 For stabs prelink tried to match GDB behavior. For N_FUN, it needs to differentiate between function start and
347 function address which are both encoded with this type, the rest of types either always need relocating or never. And
348 similarly to DWARF 2 handling, it needs to reject unknown types.

349 The relocation code in prelink is a little bit more generic than what is described above, as it is used also by other parts
350 of prelink, when growing sections in a middle of the shared library during REL to RELA conversion. All adjustment
351 functions get passed both the offset it should add to virtual addresses and a start address. Adjustment is only done if
352 the old virtual address was bigger or equal than the start address.

7 REL to RELA conversion

353 On architectures which normally use the REL format for relocations instead of RELA (IA-32, ARM and MIPS), if
354 certain relocation types use the memory r_offset points to during relocation, prelink has to either convert them to
355 a different relocation type which doesn't use the memory value, or the whole .rel.dyn section needs to be converted
356 to RELA format. Let's describe it on an example on IA-32 architecture:

```
357 $ cat > test1.c <<EOF
358 extern int i[4];
359 int *j = i + 2;
360 EOF
```

```

361 $ cat > test2.c <<EOF
362 int i[4];
363 EOF
364 $ gcc -fPIC -shared -o test2.so test2.c
365 $ gcc -fPIC -shared -o test1.so test1.c ./test2.so
366 $ readelf -l test1.so | grep LOAD | head -1
367     LOAD          0x00000000 0x0000000000 0x0000000000 0x002b8 0x002b8 R E 0x1000
368 $ readelf -l test2.so | grep LOAD | head -1
369     LOAD          0x00000000 0x0000000000 0x0000000000 0x00244 0x00244 R E 0x1000
370 $ readelf -r test1.so
371
372 Relocation section '.rel.dyn' at offset 0x2b0 contains 1 entries:
373   Offset     Info      Type            Sym.Value  Sym. Name
374 000012b8  00000d01 R_386_32           00000000    i
375 $ objdump -s -j .data test1.so
376
377 test1.so:      file format elf32-i386
378
379 Contents of section .data:
380 12b8 08000000
381 $ readelf -s test2.so | grep i\$
382     11: 000012a8  16 OBJECT  GLOBAL DEFAULT    8 i
383 $ prelink -N ./test1.so ./test2.so
384 $ readelf -l test1.so | grep LOAD | head -1
385     LOAD          0x00000000 0x04dba000 0x04dba000 0x002bc 0x002bc R E 0x1000
386 $ readelf -l test2.so | grep LOAD | head -1
387     LOAD          0x00000000 0x04db6000 0x04db6000 0x00244 0x00244 R E 0x1000
388 $ readelf -r test1.so
389
390 Relocation section '.rel.dyn' at offset 0x2b0 contains 1 entries:
391   Offset     Info      Type            Sym.Value  Sym. Name + Addend
392 04dbb2bc  00000d01 R_386_32           00000000    i + 8
393 $ objdump -s -j .data test1.so
394
395 test1.so:      file format elf32-i386
396
397 Contents of section .data:
398 4dbb2bc b072db04
399 $ readelf -s test2.so | grep i\$
400     11: 04db72a8  16 OBJECT  GLOBAL DEFAULT    8 i

```

Listing 2: REL to RELA conversion example

401 This relocation is against *i* + 8, where the addend is stored at the memory location pointed by *r_offset*. Prelink
 402 assigned base address 0x4dba000 to test1.so and 0x4db6000 to test2.so. Prelink above converted the REL
 403 section in test1.so to RELA, but let's assume it did not. All output containing 2bc above would change to 2b8
 404 (that changed above only because .rel.dyn section grew up by 4 bytes during the conversion to RELA format), the
 405 rest would stay unchanged. When some program linked against test1.so was prelinked, the (only) relocation in
 406 test1.so would not be used and *j* would contain the right value, 0x4db72b0 (address of *i* + 8; note that IA-32 is little
 407 endian, so the values in .data section are harder to read for a human). Now, let's assume one of the shared libraries
 408 the executable is linked against is upgraded. This means prelink information cannot be used, as it is out of date. Let's
 409 assume it was a library other than test2.so. Normal relocation processing for test1.so needs to happen. Standard
 410 R_386_32 calculation is *S* + *A*, in this case 0x4db72a8 + 0x4db72b0 = 0x9b6e558 and *j* contains wrong value. Either
 411 test2.so could change and now the *i* variable would have different address, or some other shared library linked to
 412 the executable could overload symbol *i*. Without additional information the dynamic linker cannot find out the addend
 413 is 8.

414 The original value of a symbol could perhaps be stored in some special allocated section and the dynamic linker could
 415 do some magic to locate it, but it would mean standard relocation handling code in the dynamic linker cannot be used
 416 for relocation processing of prelinked shared libraries where prelinking information cannot be used. So prelink in
 417 this case converts the whole .rel.dyn section into the RELA format, the addend is stored in *r_addend* field and when

418 doing relocation processing, it really doesn't matter what value is at the memory location pointed by `r_offset`. The
419 disadvantage of this is that the relocation section grew by 50%. If prelinking information can be used, it shouldn't
420 matter much, since the section is never loaded at runtime because it is not accessed. If prelinking cannot be used,
421 whether because it is out of date or because the shared library has been loaded by `dlopen`, it will increase memory
422 footprint, but it is read-only memory which is typically not used after startup and can be discarded as it is backed out
423 by the file containing the shared library.

424 At least on IA-32, REL to RELA conversion is not always necessary. If `R_386_32` added is originally 0, `prelink`
425 can instead change its type to `R_386_GLOB_DAT`, which is a similar dynamic relocation, but calculated as `S` instead of
426 `S + A`. There is no similar conversion for `R_386_PC32` possible though, on the other side this relocation type should
427 never appear in position independent shared libraries, only in position dependent code. On ARM, the situation is the
428 same, just using different relocation names (`R_ARM_32`, `R_ARM_GLOB_DAT` and `R_ARM_PC24`).

429 The `.rel.plt` section doesn't have to be converted to RELA format on either of these architectures, if the conversion is
430 needed, all other `.rel.*` allocated sections, which have to be adjacent as they are pointed to by `DT_REL` and `DT_RELSZ`
431 dynamic tags, have to be converted together. The conversion itself is fairly easy, some architecture specific code just has
432 to fetch the original addend from memory pointed by the relocation and store it into `r_addend` field (or clear `r_addend`
433 if the particular relocation type never uses the addend). The main problem is that when the conversion happens, the
434 `.rel.dyn` section grows by 50% and there needs to be room for that in the read-only loadable segment of the shared
435 library.

436 In shared libraries it is always possible to grow the first read-only `PT_LOAD` segment by adding the additional data at the
437 beginning of the read-only segment, as the shared library is relocatable. `Prelink` can relocate the whole shared library
438 to a higher address than it has assigned for it. The file offsets of all sections and the section header table file offset
439 need to be increased, but the `ELF` header and program headers need to stay at the beginning of the file. The relocation
440 section can then be moved to the newly created space between the end of the program header table and the first section.

441 Moving the section from the old location to the newly created space would leave often very big gap in virtual address
442 space as well as in the file at the old location of the relocation section. Fortunately the linker typically puts special
443 `ELF` sections including allocated relocation section before the code section and other read-only sections under user's
444 control. These special sections are intended for dynamic linking only. Their addresses are stored just in the `.dynamic`
445 section and `prelink` can easily adjust them there. There is no need for a shared library to store address of one of the
446 special sections into its code or data sections and existing linkers in fact don't create such references. When growing
447 the relocation section, `prelink` checks whether all sections before the relocation section are special¹² and if they are,
448 just moves them to lower addresses, so that the newly created space is right above the relocation section. The advantage
449 is that instead of moving all sections by the size of the new relocation section they can be adjusted ideally just by the
450 difference between old and new relocation section size.

451 There are two factors which can increase the necessary adjustment of all higher sections. The first is required section
452 alignment of any allocated section above the relocation section. `Prelink` needs to find the highest section alignment
453 among those sections and increase the adjustment from the difference between old and new relocation section up to the
454 next multiple of that alignment.

455 The second factor is only relevant to shared libraries where linker optimized the data segment placement. Traditionally
456 linker assigned the end address of the read-only segment plus the architecture's maximum `ELF` page size as the start
457 address of the read-write segment. While this created smallest file sizes of the shared libraries, it often wasted one
458 page in the read-write segment because of partial pages. When linker optimizes such that less space is wasted in partial
459 pages, the distance between read-only and read-write segments can be smaller than architecture specific maximum `ELF`
460 page size. `Prelink` has to take this into account, so that when adjusting the sections the read-only and read-write
461 segment don't end up on the same page. Unfortunately `prelink` cannot increase or decrease the distance between the
462 read-only and read-write segments, since it is possible that the shared library has relative addresses of any allocated
463 code, data or `.bss` sections stored in its sections without any relocations which would allow `prelink` to change them.
464 `Prelink` has to move all sections starting with the first allocated `SHT_PROGBITS` section other than `.interp` up to the
465 last allocated `SHT_PROGBITS` or `SHT_NOBITS` section as a block and thus needs to increase the adjustment in steps of
466 the highest section alignment as many times times as needed so that the segments end up in different pages. Below are
467 3 examples:

¹²As special sections `prelink` considers sections with `SHT_NOTE`, `SHT_HASH`, `SHT_DYNSYM`, `SHT_STRTAB`, `SHT_GNU_verdef`,
`SHT_GNU_verneed`, `SHT_GNU_versym`, `SHT_REL` or `SHT_RELA` type or the `.interp` section.

```

468 $ cat > test1.c <<EOF
469 int i[2] __attribute__((aligned (32)));
470 #define J1(N) int *j##N = &i[1];
471 #define J2(N) J1(N##0) J1(N##1) J1(N##2) J1(N##3) J1(N##4)
472 #define J3(N) J2(N##0) J2(N##1) J2(N##2) J2(N##3) J2(N##4)
473 #define J4(N) J3(N##0) J3(N##1) J3(N##2) J3(N##3) J3(N##4)
474 J4(0) J4(1) J3(2) J3(3) J1(4)
475 const int l[256] = { [10] = 1 };
476 /* Put a zero sized section at the end of read-only segment,
477 so that the end address of the segment is printed. */
478 asm (.section ro_seg_end, \"a\"; .previous");
479 EOF
480 $ gcc -shared -O2 -nostdlib -fpic -o test1.so test1.c
481 $ readelf -S test1.so | grep '^ \['
482 [Nr] Name           Type      Addr     Off   Size    ES Flg Lk Inf Al
483 [ 0] NULL           NULL      00000000 000000 000000 00 0 0 0 0
484 [ 1] .hash          HASH      000000b4 0000b4 000930 04 A 2 0 0 4
485 [ 2] .dynsym        DYNSYM   000009e4 0009e4 001430 10 A 3 d 4
486 [ 3] .dynstr        STRTAB   00001e14 001e14 000735 00 A 0 0 0 1
487 [ 4] .rel.dyn       REL      0000254c 00254c 000968 08 A 2 0 0 4
488 [ 5] .text          PROGBITS 00002eb4 002eb4 000000 00 AX 0 0 0 4
489 [ 6] .rodata         PROGBITS 00002ec0 002ec0 000400 00 A 0 0 0 32
490 [ 7] ro_seg_end    PROGBITS 000032c0 0032c0 000000 00 A 0 0 0 1
491 [ 8] .data          PROGBITS 000042c0 0032c0 0004b4 00 WA 0 0 0 4
492 [ 9] .dynamic       DYNAMIC  00004774 003774 000070 08 WA 3 0 0 4
493 [10] .got           PROGBITS 000047e4 0037e4 00000c 04 WA 0 0 0 4
494 [11] .bss           NOBITS   00004800 003800 000008 00 WA 0 0 0 32
495 [12] .comment       PROGBITS 00000000 003800 000033 00 0 0 0 1
496 [13] .shstrtab      STRTAB   00000000 003833 000075 00 0 0 0 1
497 [14] .symtab        SYMTAB   00000000 003b28 001470 10 15 11 4
498 [15] .strtab        STRTAB   00000000 004f98 000742 00 0 0 0 1
499 $ readelf -l test1.so | grep LOAD
500 LOAD      0x000000 0x00000000 0x00000000 0x032c0 0x032c0 R E 0x1000
501 LOAD      0x0032c0 0x000042c0 0x000042c0 0x00530 0x00548 RW 0x1000
502 $ prelink -N ./test1.so
503 $ readelf -l test1.so | grep LOAD
504 LOAD      0x000000 0x02000000 0x02000000 0x03780 0x03780 R E 0x1000
505 LOAD      0x003780 0x02004780 0x02004780 0x00530 0x00548 RW 0x1000
506 $ readelf -S test1.so | grep '^ \['
507 [Nr] Name           Type      Addr     Off   Size    ES Flg Lk Inf Al
508 [ 0] NULL           NULL      00000000 000000 000000 00 0 0 0 0
509 [ 1] .hash          HASH      020000b4 0000b4 000930 04 A 2 0 0 4
510 [ 2] .dynsym        DYNSYM   020009e4 0009e4 001430 10 A 3 d 4
511 [ 3] .dynstr        STRTAB   02001e14 001e14 000735 00 A 0 0 0 1
512 [ 4] .rel.dyn       RELA     0200254c 00254c 000elc 0c A 2 0 0 4
513 [ 5] .text          PROGBITS 02003374 003374 000000 00 AX 0 0 0 4
514 [ 6] .rodata         PROGBITS 02003380 003380 000400 00 A 0 0 0 32
515 [ 7] ro_seg_end    PROGBITS 02003780 003780 000000 00 A 0 0 0 1
516 [ 8] .data          PROGBITS 02004780 003780 0004b4 00 WA 0 0 0 4
517 [ 9] .dynamic       DYNAMIC  02004c34 003c34 000070 08 WA 3 0 0 4
518 [10] .got           PROGBITS 02004ca4 003ca4 00000c 04 WA 0 0 0 4
519 [11] .bss           NOBITS   02004cc0 003cc0 000008 00 WA 0 0 0 32
520 [12] .comment       PROGBITS 00000000 003cc0 000033 00 0 0 0 1
521 [13] .gnu.liblist   GNU_LIBLIST 00000000 003cf3 000000 14 14 0 0 4
522 [14] .gnu.libstr    STRTAB   00000000 003cf3 000000 00 0 0 0 1
523 [15] .gnu.prelink_undo PROGBITS 00000000 003cf4 00030c 01 0 0 0 4
524 [16] .shstrtab      STRTAB   00000000 004003 0000a0 00 0 0 0 1
525 [17] .symtab        SYMTAB   00000000 0043a0 001470 10 18 11 4
526 [18] .strtab        STRTAB   00000000 005810 000742 00 0 0 0 1

```

Listing 3: Growing read-only segment with segment distance one page

⁵²⁷ In this example the read-write segment starts at address 0x42c0, which is one page above the end of read-only segment.

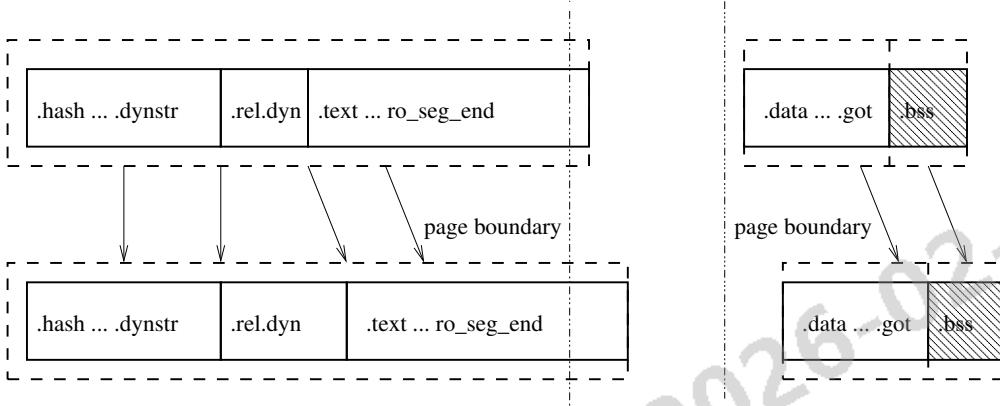


Figure 1: Growing read-only segment with segment distance one page

528 Prelink needs to grow the read-only PT_LOAD segment by 50% of .rel.dyn size, i.e. 0x4b4 bytes. Prelink just
 529 needs to round that up for the highest alignment (32 bytes required by .rodata or .bss sections) and moves all
 530 sections above .rel.dyn by 0x4c0 bytes.

```

531 $ cat > test2.c <<EOF
532 int i[2] __attribute__((aligned (32)));
533 #define J1(N) int *j##N = &i[1];
534 #define J2(N) J1(N##0) J1(N##1) J1(N##2) J1(N##3) J1(N##4)
535 #define J3(N) J2(N##0) J2(N##1) J2(N##2) J2(N##3) J2(N##4)
536 #define J4(N) J3(N##0) J3(N##1) J3(N##2) J3(N##3) J3(N##4)
537 J4(0) J4(1) J3(2) J3(3) J1(4)
538 const int l[256] = { [10] = 1 };
539 int k[670];
540 asm (" .section ro_seg_end, \"a\"; .previous");
541 EOF
542 $ gcc -shared -O2 -nostdlib -fpic -o test2.so test2.c
543 $ readelf -S test2.so | grep '^ \['
544 [Nr] Name          Type        Addr      Off      Size     ES Flg Lk Inf Al
545 [ 0]              NULL       00000000 0000000 0000000 00      0 0 0 0
546 [ 1] .hash         HASH       000000b4 0000b4 000934 04 A 2 0 4
547 [ 2] .dynsym      DYNSYM    000009e8 0009e8 001440 10 A 3 d 4
548 [ 3] .dynstr      STRTAB   00001e28 001e28 000737 00 A 0 0 1
549 [ 4] .rel.dyn     REL       00002560 002560 000968 08 A 2 0 4
550 [ 5] .text         PROGBITS 00002ec8 002ec8 0000000 00 AX 0 0 4
551 [ 6] .rodata       PROGBITS 00002ee0 002ee0 000400 00 A 0 0 32
552 [ 7] ro_seg_end   PROGBITS 000032e0 0032e0 0000000 00 A 0 0 1
553 [ 8] .data         PROGBITS 00004000 004000 0004b4 00 WA 0 0 4
554 [ 9] .dynamic     DYNAMIC   000044b4 0044b4 000070 08 WA 3 0 4
555 [10] .got         PROGBITS 00004524 004524 00000c 04 WA 0 0 4
556 [11] .bss          NOBITS   00004540 004540 000a88 00 WA 0 0 32
557 [12] .comment     PROGBITS 00000000 004540 000033 00 0 0 1
558 [13] .shstrtab    STRTAB   00000000 004573 000075 00 0 0 1
559 [14] .symtab      SYMTAB   00000000 004868 001480 10 15 11 4
560 [15] .strtab      STRTAB   00000000 005ce8 000744 00 0 0 1
561 $ readelf -l test2.so | grep LOAD
562 LOAD      0x0000000 0x000000000 0x000000000 0x032e0 0x032e0 R E 0x1000
563 LOAD      0x0040000 0x000040000 0x000040000 0x00530 0x00fc8 RW 0x1000
564 $ prelink -N ./test2.so
565 $ readelf -l test2.so | grep LOAD
566 LOAD      0x0000000 0x020000000 0x020000000 0x037a0 0x037a0 R E 0x1000
567 LOAD      0x0044c00 0x020044c00 0x020044c00 0x00530 0x00fc8 RW 0x1000
568 $ readelf -S test2.so | grep '^ \['
569 [Nr] Name          Type        Addr      Off      Size     ES Flg Lk Inf Al
570 [ 0]              NULL       00000000 0000000 0000000 00      0 0 0 0
571 [ 1] .hash         HASH       020000b4 0000b4 000934 04 A 2 0 4

```

```

572 [ 2] .dynsym      DYNSYM          020009e8 0009e8 001440 10 A 3 d 4
573 [ 3] .dynstr      STRTAB          02001e28 001e28 000737 00 A 0 0 1
574 [ 4] .rel.dyn     RELA           02002560 002560 000e1c 0c A 2 0 4
575 [ 5] .text         PROGBITS        02003388 003388 000000 00 AX 0 0 4
576 [ 6] .rodata       PROGBITS        020033a0 0033a0 000400 00 A 0 0 32
577 [ 7] ro_seg_end   PROGBITS        020037a0 0037a0 000000 00 A 0 0 1
578 [ 8] .data         PROGBITS        020044c0 0044c0 0004b4 00 WA 0 0 4
579 [ 9] .dynamic      DYNAMIC         02004974 004974 000070 08 WA 3 0 4
580 [10] .got          PROGBITS        020049e4 0049e4 00000c 04 WA 0 0 4
581 [11] .bss          NOBITS          02004a00 004a00 000a88 00 WA 0 0 32
582 [12] .comment      PROGBITS        00000000 004a00 000033 00 0 0 1
583 [13] .gnu.liblist   GNU_LIBLIST    00000000 004a33 000000 14 14 0 4
584 [14] .gnu.libstr    STRTAB          00000000 004a33 000000 00 0 0 1
585 [15] .gnu.prelink_undo PROGBITS    00000000 004a34 00030c 01 0 0 4
586 [16] .shstrtab     STRTAB          00000000 004d43 0000a0 00 0 0 1
587 [17] .symtab       SYMTAB          00000000 0050e0 001480 10 18 11 4
588 [18] .strtab       STRTAB          00000000 006560 000744 00 0 0 1

```

Listing 4: Growing read-only segment not requiring additional padding

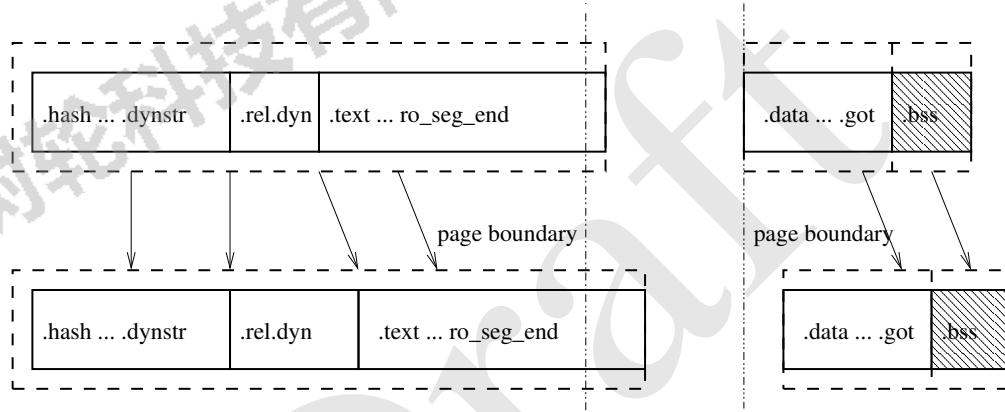


Figure 2: Growing read-only segment not requiring additional padding

589 In the second example `prelink` can grow by just `0x4c0` bytes as well, even though the distance between read-write
590 and read-only segment is just `0xd20` bytes. With this distance, hypothetical adjustment by any size less than `0xd21`
591 bytes (modulo 4096) would need just rounding up to the next multiple of 32 bytes, while adjustments from `0xd21` up
592 to `0xfe0` would require adjustments in multiples of 4096 bytes.

```

593 $ cat > test3.c <<EOF
594 int i[2] __attribute__((aligned (32)));
595 #define J1(N) int *j##N = &i[1];
596 #define J2(N) J1(N##0) J1(N##1) J1(N##2) J1(N##3) J1(N##4)
597 #define J3(N) J2(N##0) J2(N##1) J2(N##2) J2(N##3) J2(N##4)
598 #define J4(N) J3(N##0) J3(N##1) J3(N##2) J3(N##3) J3(N##4)
599 J4(0) J4(1) J3(2) J3(3) J1(4)
600 int k[670];
601 asm (" .section ro_seg_end, \"a\"; .previous");
602 EOF
603 $ gcc -shared -O2 -nostdlib -fpic -o test3.so test3.c
604 $ readelf -S test3.so | grep '^ \['
605 [Nr] Name             Type            Addr     Off      Size     ES Flg Lk Inf Al
606 [ 0]                 NULL           00000000 000000 000000 00 0 0 0 0
607 [ 1] .hash            HASH           000000b4 0000b4 00092c 04 A 2 0 4
608 [ 2] .dynsym          DYNSYM         000009e0 0009e0 001420 10 A 3 c 4
609 [ 3] .dynstr          STRTAB         00001e00 001e00 000735 00 A 0 0 1
610 [ 4] .rel.dyn         REL            00002538 002538 000968 08 A 2 0 4

```

```

611 [ 5] .text           PROGBITS    00002ea0 002ea0 000000 00 AX 0 0 4
612 [ 6] ro_seg_end     PROGBITS    00002ea0 002ea0 000000 00 A 0 0 1
613 [ 7] .data           PROGBITS    00003000 003000 0004b4 00 WA 0 0 4
614 [ 8] .dynamic        DYNAMIC     000034b4 0034b4 000070 08 WA 3 0 4
615 [ 9] .got            PROGBITS    00003524 003524 00000c 04 WA 0 0 4
616 [10] .bss             NOBITS      00003540 003540 000a88 00 WA 0 0 32
617 [11] .comment         PROGBITS    00000000 003540 000033 00 0 0 1
618 [12] .shstrtab        STRTAB      00000000 003573 00006d 00 0 0 1
619 [13] .symtab          SYMTAB      00000000 003838 001460 10 14 10 4
620 [14] .strtab          STRTAB      00000000 004c98 000742 00 0 0 1
621 $ readelf -l test3.so | grep LOAD
622 LOAD      0x000000 0x00000000 0x00000000 0x02ea0 0x02ea0 R E 0x1000
623 LOAD      0x003000 0x00003000 0x00003000 0x00530 0x00fc8 RW 0x1000
624 $ prelink -N ./test3.so
625 $ readelf -l test3.so | grep LOAD
626 LOAD      0x000000 0x02000000 0x02000000 0x03ea0 0x03ea0 R E 0x1000
627 LOAD      0x004000 0x02004000 0x02004000 0x00530 0x00fc8 RW 0x1000
628 $ readelf -S test3.so | grep '^ \['
629 [Nr] Name           Type        Addr   Off    Size   ES Flg Lk Inf Al
630 [ 0]                NULL       00000000 000000 000000 00 0 0 0 0
631 [ 1] .hash           HASH       020000b4 0000b4 00092c 04 A 2 0 4
632 [ 2] .dynsym         DYNSYM     020009e0 0009e0 001420 10 A 3 c 4
633 [ 3] .dynstr         STRTAB     02001e00 001e00 000735 00 A 0 0 1
634 [ 4] .rel.dyn        RELA      02002538 002538 000e1c 0c A 2 0 4
635 [ 5] .text           PROGBITS  02003ea0 003ea0 000000 00 AX 0 0 4
636 [ 6] ro_seg_end     PROGBITS  02003ea0 003ea0 000000 00 A 0 0 1
637 [ 7] .data           PROGBITS  02004000 004000 0004b4 00 WA 0 0 4
638 [ 8] .dynamic        DYNAMIC    020044b4 0044b4 000070 08 WA 3 0 4
639 [ 9] .got            PROGBITS  02004524 004524 00000c 04 WA 0 0 4
640 [10] .bss             NOBITS    02004540 004540 000a88 00 WA 0 0 32
641 [11] .comment         PROGBITS  00000000 004540 000033 00 0 0 1
642 [12] .gnu.liblist    GNU_LIBLIST 00000000 004573 000000 14 13 0 4
643 [13] .gnu.libstr     STRTAB    00000000 004573 000000 00 0 0 1
644 [14] .gnu.prelink_undo PROGBITS  00000000 004574 0002e4 01 0 0 4
645 [15] .shstrtab        STRTAB    00000000 00485b 000098 00 0 0 1
646 [16] .symtab          SYMTAB    00000000 004bc8 001460 10 17 10 4
647 [17] .strtab          STRTAB    00000000 006028 000742 00 0 0 1

```

Listing 5: Growing read-only segment if page padding needed

648 In the last example the distance between PT_LOAD segments is very small, just 0x160 bytes and the adjustment had to
649 be done by 4096 bytes.

8 Conflicts

650 As said earlier, if symbol lookup of some symbol in particular shared library results in different values when that
651 shared library's natural search scope is used and when using search scope of the application the DSO is used in, this is
652 considered a *conflict*. Here is an example of a conflict on IA-32:

```

653 $ cat > test1.c <<EOF
654 int i;
655 int *j = &i;
656 int *foo (void) { return &i; }
657 EOF
658 $ cat > test2.c <<EOF
659 int i;
660 int *k = &i;
661 int *bar (void) { return &i; }
662 EOF
663 $ cat > test.c <<EOF

```

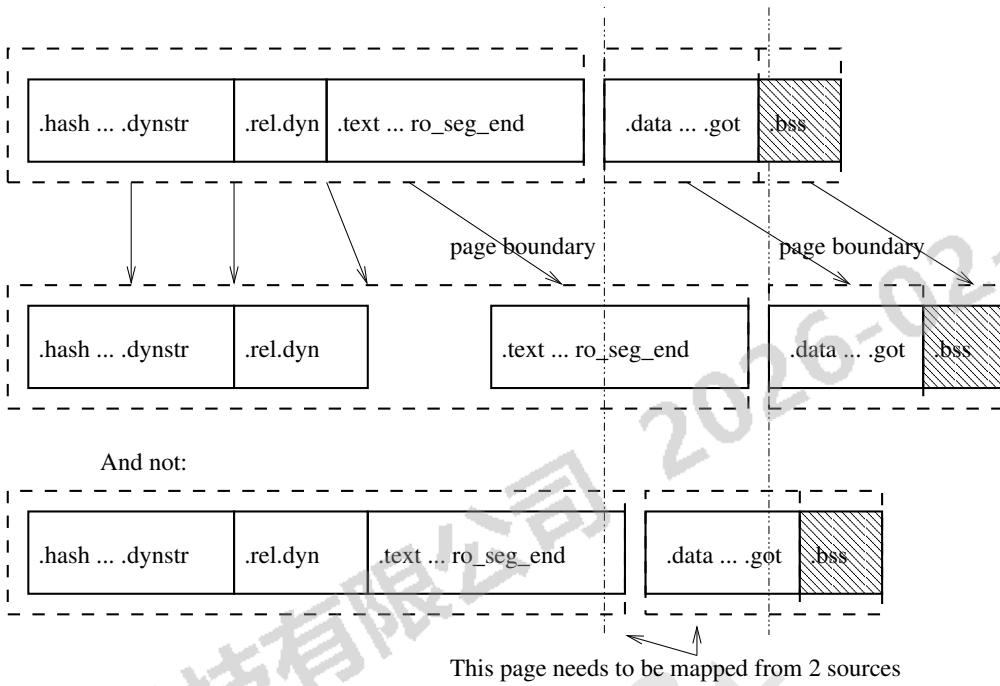


Figure 3: Growing read-only segment if page padding needed

```

664 #include <stdio.h>
665 extern int i, *j, *k, *foo (void), bar (void);
666 int main (void)
667 {
668 #ifdef PRINT_I
669     printf ("%p\n", &i);
670 #endif
671     printf ("%p %p %p %p\n", j, k, foo (), bar ());
672 }
673 EOF
674 $ gcc -nostdlib -shared -fpic -s -o test1.so test1.c
675 $ gcc -nostdlib -shared -fpic -o test2.so test2.c ./test1.so
676 $ gcc -o test test.c ./test2.so ./test1.so
677 $ ./test
678 0x16137c 0x16137c 0x16137c 0x16137c
679 $ readelf -r ./test1.so
680
681 Relocation section '.rel.dyn' at offset 0x2bc contains 2 entries:
682   Offset     Info      Type            Sym.Value  Sym. Name
683 000012e4  00000d01 R_386_32        00001368  i
684 00001364  00000d06 R_386_GLOB_DAT  00001368  i
685 $ prelink -N ./test ./test1.so ./test2.so
686 $ LD_WARN= LD_TRACE_PRELINKING=1 LD_BIND_NOW=1 /lib/ld-linux.so.2 ./test1.so
687         ./test1.so => ./test1.so (0x04db6000, 0x00000000)
688 $ LD_WARN= LD_TRACE_PRELINKING=1 LD_BIND_NOW=1 /lib/ld-linux.so.2 ./test2.so
689         ./test2.so => ./test2.so (0x04dba000, 0x00000000)
690         ./test1.so => ./test1.so (0x04db6000, 0x00000000)
691 $ LD_WARN= LD_TRACE_PRELINKING=1 LD_BIND_NOW=1 /lib/ld-linux.so.2 ./test \
692 | sed 's/^[[[:space:]]]*//'
693 ./test => ./test (0x08048000, 0x00000000)
694 ./test2.so => ./test2.so (0x04dba000, 0x00000000)
695 ./test1.so => ./test1.so (0x04db6000, 0x00000000)
696 libc.so.6 => /lib/tls/libc.so.6 (0x00b22000, 0x00000000) TLS(0x1, 0x00000028)
697 /lib/ld-linux.so.2 => /lib/ld-linux.so.2 (0x00b0a000, 0x00000000)
698 $ readelf -S ./test1.so | grep '\.data||\.got'
699 [ 6] .data          PROGBITS    04db72e4 0002e4 000004 00 WA 0 0 4
700 [ 8] .got          PROGBITS    04db7358 000358 000010 04 WA 0 0 4

```

```

701 $ readelf -r ./test1.so
702
703 Relocation section '.rel.dyn' at offset 0x2bc contains 2 entries:
704   Offset     Info      Type            Sym.Value  Sym. Name
705   04db72e4  00000d06 R_386_GLOB_DAT    04db7368  i
706   04db7364  00000d06 R_386_GLOB_DAT    04db7368  i
707 $ objdump -s -j .got -j .data test1.so
708
709 test1.so:      file format elf32-i386
710
711 Contents of section .data:
712   4db72e4 6873db04          hs..
713 Contents of section .got:
714   4db7358 e8120000 00000000 00000000 6873db04 .....hs..
715 $ readelf -r ./test | sed '/\.\.gnu\.conflict/,\$!d'
716 Relocation section '.gnu.conflict' at offset 0x7ac contains 18 entries:
717   Offset     Info      Type            Sym.Value  Sym. Name + Addend
718   04db72e4  00000001 R_386_32          04dbb37c
719   04db7364  00000001 R_386_32          04dbb37c
720   00c56874  00000001 R_386_32          fffffff0
721   00c56878  00000001 R_386_32          00000001
722   00c568bc  00000001 R_386_32          fffffff4
723   00c56900  00000001 R_386_32          ffffffec
724   00c56948  00000001 R_386_32          ffffffdc
725   00c5695c  00000001 R_386_32          ffffffe0
726   00c56980  00000001 R_386_32          ffffff8
727   00c56988  00000001 R_386_32          ffffffe4
728   00c569a4  00000001 R_386_32          ffffffd8
729   00c569c4  00000001 R_386_32          ffffffe8
730   00c569d8  00000001 R_386_32          080485b8
731   00b1f510  00000007 R_386_JUMP_SLOT  00b91460
732   00b1f514  00000007 R_386_JUMP_SLOT  00b91080
733   00b1f518  00000007 R_386_JUMP_SLOT  00b91750
734   00b1f51c  00000007 R_386_JUMP_SLOT  00b912c0
735   00b1f520  00000007 R_386_JUMP_SLOT  00b91200
736 $ ./test
737 0x4dbb37c 0x4dbb37c 0x4dbb37c 0x4dbb37c

```

Listing 6: Conflict example

738 In the example, among some conflicts caused by the dynamic linker and the C library,¹³ there is a conflict for the
 739 symbol *i* in *test1.so* shared library. *test1.so* has just itself in its natural symbol lookup scope (as proved by

```

740 LD_WARN= LD_TRACE_PRELINKING=1 LD_BIND_NOW=1 /lib/ld-linux.so.2 ./test1.so

```

741 command output), so when looking up symbol *i* in this scope the definition in *test1.so* is chosen. *test1.so* has
 742 two relocations against the symbol *i*, one *R_386_32* against *.data* section and one *R_386_GLOB_DAT* against *.got*
 743 section. When prelinking *test1.so* library, the dynamic linker stores the address of *i* (0x4db7368) into both locations
 744 (at offsets 0x4db72e4 and 0x4db7364). The global symbol search scope in *test* executable contains the executable
 745 itself, *test2.so* and *test1.so* libraries, *libc.so.6* and the dynamic linker in the listed order. When doing symbol
 746 lookup for symbol *i* in *test1.so* when doing relocation processing of the whole executable, address of *i* in *test2.so*
 747 is returned as that symbol comes earlier in the global search scope. So, when none of the libraries nor the executable
 748 is prelinked, the program prints 4 identical addresses. If prelink didn't create conflict fixups for the two relocations
 749 against the symbol *i* in *test1.so*, prelinked executable (which bypasses normal relocation processing on startup)
 750 would print instead of the desired

```

751 0x4dbb37c 0x4dbb37c 0x4dbb37c 0x4dbb37c

```

¹³Particularly in the example, the 5 *R_386_JUMP_SLOT* fixups are PLT slots in the dynamic linker for memory allocator functions resolving to C library functions instead of dynamic linker's own trivial implementation. First 10 *R_386_32* fixups at offsets 0xc56874 to 0xc569c4 are Thread Local Storage fixups in the C library and the fixup at 0xc569d8 is for *_IO_stdin_used* weak undefined symbol in the C library, resolving to a symbol with the same name in the executable.

752 different addresses,

753 0x4db7368 0x4dbb37c 0x4db7368 0x4dbb37c

754 That is a functionality change that prelink cannot be permitted to make, so instead it fixes up the two locations by
755 storing the desired value in there. In this case prelink really cannot avoid that - test1.so shared library could
756 be also used without test2.so in some other executable's symbol search scope. Or there could be some executable
757 linked with:

758 \$ gcc -o test2 test.c ./test1.so ./test2.so

Listing 7: Conflict example with swapped order of libraries

759 where *i* lookup in test1.so and test2.so is supposed to resolve to *i* in test1.so.

760 Now consider what happens if the executable is linked with -DPRINT_I:

```
761 $ gcc -DPRINT_I -o test3 test.c ./test2.so ./test1.so
762 $ ./test3
763 0x804972c
764 0x804972c 0x804972c 0x804972c 0x804972c
765 $ prelink -N ./test3 ./test1.so ./test2.so
766 $ readelf -S ./test2.so | grep '\.data\|\.got'
767 [ 6 ] .data           PROGBITS          04dbb2f0 0002f0 000004 00 WA 0 0 4
768 [ 8 ] .got            PROGBITS          04dbb36c 00036c 000010 04 WA 0 0 4
769 $ readelf -r ./test2.so
770
771 Relocation section '.rel.dyn' at offset 0x2c8 contains 2 entries:
772 Offset     Info      Type      Sym.Value  Sym. Name
773 04dbb2f0  00000d06 R_386_GLOB_DAT    04dbb37c  i
774 04dbb378  00000d06 R_386_GLOB_DAT    04dbb37c  i
775 $ objdump -s -j .got -j .data test2.so
776
777 test2.so:      file format elf32-i386
778
779 Contents of section .data:
780 4dbb2f0 7cb3db04          ...
781 Contents of section .got:
782 4dbb36c f4120000 00000000 00000000 7cb3db04  .....
783 $ readelf -r ./test3
784
785 Relocation section '.rel.dyn' at offset 0x370 contains 4 entries:
786 Offset     Info      Type      Sym.Value  Sym. Name
787 08049720  00000e06 R_386_GLOB_DAT    00000000  __gmon_start__
788 08049724  00000105 R_386_COPY       08049724  j
789 08049728  00000305 R_386_COPY       08049728  k
790 0804972c  00000405 R_386_COPY       0804972c  i
791
792 Relocation section '.rel.plt' at offset 0x390 contains 4 entries:
793 Offset     Info      Type      Sym.Value  Sym. Name
794 08049710  00000607 R_386_JUMP_SLOT  080483d8  __libc_start_main
795 08049714  00000707 R_386_JUMP_SLOT  080483e8  printf
796 08049718  00000807 R_386_JUMP_SLOT  080483f8  foo
797 0804971c  00000c07 R_386_JUMP_SLOT  08048408  bar
798
799 Relocation section '.gnu.conflict' at offset 0x7f0 contains 20 entries:
800 Offset     Info      Type      Sym.Value  Sym. Name + Addend
801 04dbb2f0  00000001 R_386_32          0804972c
```

```

802 04dbb378 00000001 R_386_32          0804972c
803 04db72e4 00000001 R_386_32          0804972c
804 04db7364 00000001 R_386_32          0804972c
805 00c56874 00000001 R_386_32          ffffff0
806 00c56878 00000001 R_386_32          00000001
807 00c568bc 00000001 R_386_32          ffffff4
808 00c56900 00000001 R_386_32          ffffffec
809 00c56948 00000001 R_386_32          ffffffdc
810 00c5695c 00000001 R_386_32          ffffffe0
811 00c56980 00000001 R_386_32          ffffff8
812 00c56988 00000001 R_386_32          ffffffe4
813 00c569a4 00000001 R_386_32          ffffffd8
814 00c569c4 00000001 R_386_32          ffffffe8
815 00c569d8 00000001 R_386_32          080485f0
816 00b1f510 00000007 R_386_JUMP_SLOT   00b91460
817 00b1f514 00000007 R_386_JUMP_SLOT   00b91080
818 00b1f518 00000007 R_386_JUMP_SLOT   00b91750
819 00b1f51c 00000007 R_386_JUMP_SLOT   00b912c0
820 00b1f520 00000007 R_386_JUMP_SLOT   00b91200
821 $ ./test3
822 0x804972c
823 0x804972c 0x804972c 0x804972c

```

Listing 8: Conflict example with COPY relocation for conflicting symbol

824 Because the executable is not compiled as position independent code and `main` function takes address of *i* variable,
 825 the object file for `test3.c` contains a `R_386_32` relocation against *i*. The linker cannot make dynamic relocations
 826 against read-only segment in the executable, so the address of *i* must be constant. This is accomplished by creating a
 827 new object *i* in the executable's `.dynbss` section and creating a dynamic `R_386_COPY` relocation for it. The relocation
 828 ensures that during startup the content of *i* object earliest in the search scope without the executable is copied to this *i*
 829 object in executable. Now, unlike `test` executable, in `test3` executable *i* lookups in both `test1.so` and `test2.so`
 830 libraries result in address of *i* in the executable (instead of `test2.so`). This means that two conflict fixups are needed
 831 again for `test1.so` (but storing `0x804972c` instead of `0x4dbb37c`) and two new fixups are needed for `test2.so`.

832 If the executable is compiled as position independent code,

```

833 $ gcc -fpic -DPRINT_I -o test4 test.c ./test2.so ./test1.so
834 $ ./test4
835 0x4dbb37c
836 0x4dbb37c 0x4dbb37c 0x4dbb37c 0x4dbb37c

```

Listing 9: Conflict example with position independent code in the executable

837 the address of *i* is stored in executable's `.got` section, which is writable and thus can have dynamic relocation against it.
 838 So the linker creates a `R_386_GLOB_DAT` relocation against the `.got` section, the symbol *i* is undefined in the executable
 839 and no copy relocations are needed. In this case, only `test1.so` will need 2 fixups, `test2.so` will not need any.

840 There are various reasons for conflicts:

- 841 • Improperly linked shared libraries. If a shared library always needs symbols from some particular shared library,
 842 it should be linked against that library, usually by adding `-LLIBNAME` to `gcc -shared` command line used
 843 during linking of the shared library. This both reduces conflict fixups in prelink and makes the library easier
 844 to load using `dlopen`, because applications don't have to remember that they have to load some other library
 845 first. The best place to record the dependency is in the shared library itself. Another reason is if the needed
 846 library uses symbol versioning for its symbols. Not linking against that library can result in malfunctioning
 847 shared library. Prelink issues a warning for such libraries - Warning: *library* has undefined non-weak

848 symbols. When linking a shared library, the `-Wl,-z,defs` option can be used to ensure there are no such
849 undefined non-weak symbols. There are exceptions, when undefined non-weak symbols in shared libraries are
850 desirable. One exception is when there are multiple shared libraries providing the same functionality, and a
851 shared library doesn't care which one is used. An example can be e.g. `libreadline.so.4`, which needs some
852 terminal handling functions, which are provided by either `libtermcap.so.2`, or `libcurses.so.5`. Another
853 exception is with plugins or other shared libraries which expect some symbols to be resolved to symbols defined
854 in the executable.

- A library overriding functionality of some other library. One example is e.g. C library and POSIX thread library. Older versions of the GNU C library did not provide cancelable entry points required by the standard. This is not needed for non-threaded applications. So only the `libpthread.so.0` shared library which provides POSIX threading support then overrode the cancellation entry points required by the standard by wrapper functions which provided the required functionality. Although most recent versions of the GNU C library handle cancellation even in entry points in `libc.so.6` (this was needed for cases when `libc.so.6` comes earlier before `libpthread.so.0` in symbol search scope and used to be worked around by non-standard handling of weak symbols in the dynamic linker), because of symbol versioning the symbols had to stay in `libpthread.so.0` as well as in `libc.so.6`. This means every program using POSIX threads on Linux will have a couple of conflict fixups because of this.
- Programs which need copy relocations. Although `prelink` will resolve the copy relocations at prelinking time, if any shared library has relocations against the symbol which needed copy relocation, all such relocations will need conflict fixups. Generally, it is better to not export variables from shared libraries in their APIs, instead provide accessor functions.
- Function pointer equality requirement for functions called from executables. When address of some global function is taken, at least C and C++ require that this pointer is the same in the whole program. Executables typically contain position dependent code, so when code in the executable takes address of some function not defined in the executable itself, that address must be link time constant. Linker accomplishes this by creating a PLT slot for the function unless there was one already and resolving to the address of PLT slot. The symbol for the function is created with `st_value` equal to address of the PLT slot, but `st_shndx` set to `SHN_UNDEF`. Such symbols are treated specially by the dynamic linker, in that PLT relocations resolve to first symbol in the global search scope after the executable, while symbol lookups for all other relocation types return the address of the symbol in the executable. Unfortunately, GNU linker doesn't differentiate between taking address of a function in an executable (especially one for which no dynamic relocation is possible in case it is in read-only segment) and just calling the function, but never taking its address. If it cleared the `st_value` field of the `SHN_UNDEF` function symbols in case nothing in the executable takes the function's address, several `prelink` conflict could disappear (`SHN_UNDEF` symbols with `st_value` set to 0 are treated always as real undefined symbols by the dynamic linker).
- COMDAT code and data in C++. C++ language has several places where it may need to emit some code or data without a clear unique compilation unit owning it. Examples include taking address of an `inline` function, local static variable in `inline` functions, virtual tables for some classes (this depends on `#pragma interface` or `#pragma implementation` presence, presence of non-`inline` non-pure-virtual member function in the class, etc.), RTTI info for them. Compilers and linkers handle these using various COMDAT schemes, e.g. GNU linker's `.gnu.linkonce*` special sections or using `SHT_GROUP`. Unfortunately, all these duplicate merging schemes work only during linking of shared libraries or executables, no duplicate removal is done across shared libraries. Shared libraries typically have relocations against their COMDAT code or data objects (otherwise they wouldn't be at least in most cases emitted at all), so if there are COMDAT duplicates across shared libraries or the executable, they lead to conflict fixups. The linker theoretically could try to merge COMDAT duplicates across shared libraries if specifically requested by the user (if a COMDAT symbol is already present in one of the dependent shared libraries and is `STB_WEAK`, the linker could skip it). Unfortunately, this only works as long as the user has full control over the dependent shared libraries, because the COMDAT symbol could be exported from them just as a side effect of their implementation (e.g. they use some class internally). When such libraries are rebuilt even with minor changes in their implementation (unfortunately with C++ shared libraries it is usually not very clear what part is exported ABI and what is not), some of those COMDAT symbols in them could go away (e.g. because suddenly they use a different class internally and the previously used class is not referenced anywhere). When COMDAT objects are not merged across shared libraries, this makes no problems, as each library which needs the COMDAT has its own copy. But with COMDAT duplicate removal between shared libraries there could suddenly be unresolved references and the shared libraries would need to be relinked. The only place where this could work safely is when a single package includes several C++ shared libraries which depend on each other. They are then shipped always together and when one changes, all others need changing too.

9 Prelink optimizations to reduce number of conflict fixups

905 Prelink can optimize out some conflict fixups if it can prove that the changes are not observable by the application
906 at runtime (opening its executable and reading it doesn't count). If there is a data object in some shared library with
907 a symbol that is overridden by a symbol in a different shared library earlier in global symbol lookup scope or in
908 the executable, then that data object is likely never referenced and it shouldn't matter what it contains. Examine the
909 following example:

```
910 $ cat > test1.c <<EOF
911 int i, j, k;
912 struct A { int *a; int *b; int *c; } x = { &i, &j, &k };
913 struct A *y = &x;
914 EOF
915 $ cat > test2.c <<EOF
916 int i, j, k;
917 struct A { int *a; int *b; int *c; } x = { &i, &j, &k };
918 struct A *z = &x;
919 EOF
920 $ cat > test.c <<EOF
921 #include <stdio.h>
922 extern struct A { int *a; int *b; int *c; } *y, *z;
923 int main (void)
924 {
925     printf ("%p: %p %p\n", y, y->a, y->b, y->c);
926     printf ("%p: %p %p\n", z, z->a, z->b, z->c);
927 }
928 EOF
929 $ gcc -nostdlib -shared -fpic -s -o test1.so test1.c
930 $ gcc -nostdlib -shared -fpic -o test2.so test2.c ./test1.so
931 $ gcc -o test test.c ./test2.so ./test1.so
932 $ ./test
933 0xaf3314: 0xaf33b0 0xaf33a8 0xaf33ac
934 0xaf3314: 0xaf33b0 0xaf33a8 0xaf33ac
```

Listing 10: C example where conflict fixups could be optimized out

935 In this example there are 3 conflict fixups pointing into the 12 byte long `x` object in `test1.so` shared library (among
936 other conflicts). And nothing in the program can poke at `x` content in `test1.so`, simply because it has to look at it
937 through `x` symbol which resolves to `test2.so`. So in this case prelink could skip those 3 conflicts. Unfortunately it
938 is not that easy:

```
939 $ cat > test3.c <<EOF
940 int i, j, k;
941 static struct A { int *a; int *b; int *c; } local = { &i, &j, &k };
942 extern struct A x;
943 struct A *y = &x;
944 struct A *y2 = &local;
945 extern struct A x __attribute__((alias ("local")));
946 EOF
947 $ cat > test4.c <<EOF
948 #include <stdio.h>
949 extern struct A { int *a; int *b; int *c; } *y, *y2, *z;
950 int main (void)
951 {
952     printf ("%p: %p %p\n", y, y->a, y->b, y->c);
953     printf ("%p: %p %p\n", y2, y2->a, y2->b, y2->c);
954     printf ("%p: %p %p\n", z, z->a, z->b, z->c);
955 }
```

```

956 EOF
957 $ gcc -fPIC -shared -o test3.so test3.c
958 $ gcc -fPIC -shared -o test4.so test2.c ./test3.so
959 $ gcc -o test4 test4.c ./test4.so ./test3.so
960 $ ./test4
961 0x65a314: 0x65a3b0 0x65a3a8 0x65a3ac
962 0xbd1328: 0x65a3b0 0x65a3a8 0x65a3ac
963 0x65a314: 0x65a3b0 0x65a3a8 0x65a3ac

```

Listing 11: Modified C example where conflict fixups cannot be removed

964 In this example, there are again 3 conflict fixups pointing into the 12 byte long `x` object in `test3.so` shared library.
 965 The fact that variable `local` is located at the same 12 bytes is totally invisible to `prelink`, as `local` is a `STB_LOCAL` symbol
 966 which doesn't show up in `.dynsym` section. But if those 3 conflict fixups are removed, then suddenly program's
 967 observable behavior changes (the last 3 addresses on second line would be different than those on first or third line).

968 Fortunately, there are at least some objects where `prelink` can be reasonably sure they will never be referenced
 969 through some local alias. Those are various compiler generated objects with well defined meaning which is `prelink`
 970 able to identify in shared libraries. The most important ones are C++ virtual tables and *RTTI* data. They are emitted
 971 as COMDAT data by the compiler, in GCC into `.gnu.linkonce.d.*` sections. Data or code in these sections can
 972 be accessed only through global symbols, otherwise linker might create unexpected results when two or more of these
 973 sections are merged together (all but one deleted). When `prelink` is checking for such data, it first checks whether the
 974 shared library in question is linked against `libstdc++.so`. If not, it is not a C++ library (or incorrectly built one) and
 975 thus it makes no sense to search any further. It looks only in `.data` section, for `STB_WEAK STT_OBJECT` symbols whose
 976 names start with certain prefixes¹⁴ and where no other symbols (in dynamic symbol table) point into the objects. If
 977 these objects are unused because there is a conflict on their symbol, all conflict fixups pointing into the virtual table or
 978 *RTTI* structure can be discarded.

979 Another possible optimization is again related to C++ virtual tables. Function addresses in them are not intended for
 980 pointer comparisons. C++ code only loads them from the virtual tables and calls through the pointer. Pointers to
 981 member functions are handled differently. As pointer equivalence is the only reason why all function pointers resolve
 982 to PLT slots in the executable even when the executable doesn't include implementation of the function (i.e. has
 983 `SHN_UNDEF` symbol with non-zero `st_value` pointing at the PLT slot in the executable), `prelink` can resolve method
 984 addresses in virtual tables to the actual method implementation. In many cases this is in the same library as the virtual
 985 table (or in one of libraries in its natural symbol lookup scope), so a conflict fixup is unnecessary. This optimization
 986 speeds up programs also after control is transferred to the application and not just the time to start up the application,
 987 although just a few cycles per method call.

988 The conflict fixup reduction is quite big on some programs. Below is statistics for `kmail` program on completely
 989 unprelinked box:

```

990 $ LD_DEBUG=statistics /usr/bin/kmail 2>&1 | sed '2,8!d;s/^ *//'
991 10621:      total startup time in dynamic loader: 240724867 clock cycles
992 10621:          time needed for relocation: 234049636 clock cycles (97.2%)
993 10621:          number of relocations: 34854
994 10621:          number of relocations from cache: 74364
995 10621:          number of relative relocations: 35351
996 10621:          time needed to load objects: 6241678 clock cycles (2.5%)
997 $ ls -l /usr/bin/kmail
998 -rwxr-xr-x  1 root      root  2149084 Oct  2 12:05 /usr/bin/kmail
999 $ ( Xvfb :3 & ) >/dev/null 2>&1 </dev/null; sleep 20
1000 $ ( DISPLAY=:3 kmail& ) >/dev/null 2>&1 </dev/null; sleep 10; killall kmail
1001 $ ( DISPLAY=:3 kmail& ) >/dev/null 2>&1 </dev/null; sleep 10
1002 $ cat /proc/'/sbin/pidof kmail'/statm
1003 4164 4164 3509 224 33 3907 655
1004 $ killall Xvfb kdeinit kmail

```

¹⁴ `_vt_` for GCC 2.95.x and 2.96-RH virtual tables, `_ZTV` for GCC 3.x virtual tables and `_ZTI` for GCC 3.x *RTTI* data.

Listing 12: Statistics for unprelinked kmail

1005 statm special file for a process contains its memory statistics. The numbers in it mean in order total number of used
1006 pages (on IA-32 Linux a page is 4KB), number of resident pages (i.e. not swapped out), number of shared pages,
1007 number of text pages, number of library pages, number of stack and other pages and number of dirty pages used by
1008 the process. Distinction between text and library pages is very rough, so those numbers aren't that much useful. Of
1009 interest are mainly first number, third number and last number.

1010 Statistics for kmail on completely prelinked box:

```
1011 $ LD_DEBUG=statistics /usr/bin/kmail 2>&1 | sed '2,8!d;s/^ *//'  
1012 14864:      total startup time in dynamic loader: 8409504 clock cycles  
1013 14864:          time needed for relocation: 3024720 clock cycles (35.9%)  
1014 14864:          number of relocations: 0  
1015 14864:          number of relocations from cache: 8961  
1016 14864:          number of relative relocations: 0  
1017 14864:          time needed to load objects: 4897336 clock cycles (58.2%)  
1018 $ ls -l /usr/bin/kmail  
1019 -rwxr-xr-x  1 root      root    2269500 Oct  2 12:05 /usr/bin/kmail  
1020 $ ( Xvfb :3 & ) >/dev/null 2>&1 </dev/null; sleep 20  
1021 $ ( DISPLAY=:3 kmail& ) >/dev/null 2>&1 </dev/null; sleep 10; killall kmail  
1022 $ ( DISPLAY=:3 kmail& ) >/dev/null 2>&1 </dev/null; sleep 10  
1023 $ cat /proc/'/sbin/pidof kmail`/statm  
1024 3803 3803 3186 249 33 3521 617  
1025 $ killall Xvfb kdeinit kmail
```

Listing 13: Statistics for prelinked kmail

1026 Statistics for kmail on completely prelinked box with C++ conflict fixup optimizations turned off:

```
1027 $ LD_DEBUG=statistics /usr/bin/kmail 2>&1 | sed '2,8!d;s/^ *//'  
1028 20645:      total startup time in dynamic loader: 9704168 clock cycles  
1029 20645:          time needed for relocation: 4734715 clock cycles (48.7%)  
1030 20645:          number of relocations: 0  
1031 20645:          number of relocations from cache: 59871  
1032 20645:          number of relative relocations: 0  
1033 20645:          time needed to load objects: 4487971 clock cycles (46.2%)  
1034 ls -l /usr/bin/kmail  
1035 -rwxr-xr-x  1 root      root    2877360 Oct  2 12:05 /usr/bin/kmail  
1036 $ ( Xvfb :3 & ) >/dev/null 2>&1 </dev/null; sleep 20  
1037 $ ( DISPLAY=:3 kmail& ) >/dev/null 2>&1 </dev/null; sleep 10; killall kmail  
1038 $ ( DISPLAY=:3 kmail& ) >/dev/null 2>&1 </dev/null; sleep 10  
1039 $ cat /proc/'/sbin/pidof kmail`/statm  
1040 3957 3957 3329 398 33 3526 628  
1041 $ killall Xvfb kdeinit kmail
```

Listing 14: Statistics for prelinked kmail without conflict fixup reduction

1042 On this application, C++ conflict fixup optimizations saved 50910 unneeded conflict fixups, speeded up startup by
1043 13.3% and decreased number of dirty pages by 11, which means the application needs 44KB less memory per-process.

10 Thread Local Storage support

1044 Thread Local Storage ([12], [13], [14]) support has been recently added to GCC, GNU binutils and GNU C Li-
1045 brary. TLS support is a set of new relocations which together with dynamic linker and POSIX thread library addi-

1046 tions provide faster and easier to use alternative to traditional POSIX thread local data API (`pthread_getspecific`,
1047 `pthread_setspecific`, `pthread_key_*`).

1048 TLS necessitated several changes to prelink. Thread Local symbols (with type `STT_TLS`) must not be relocated, as
1049 they are relative to the start of `PT_TLS` segment and thus not virtual addresses. The dynamic linker had to be enhanced
1050 so that it tells prelink at `LD_TRACE_PRELINKING` time what TLS module IDs have been assigned and what addresses
1051 relative to start of TLS block have been given to `PT_TLS` segment of each library or executable. There are 3 classes of
1052 new TLS dynamic relocations prelink is interested in (with different names on different architectures).

1053 In first class are module ID relocations, which are used for TLS Global Dynamic and Local Dynamic models (for Global
1054 Dynamic model they are supposed to resolve to module ID of the executable or shared library of particular `STT_TLS`
1055 symbol, for Local Dynamic model this resolves to module ID of the containing shared library). These relocations are
1056 hard to prelink in any useful way without moving TLS module ID assignment from the dynamic linker to prelink.
1057 Although prelink can find out what shared library will contain particular `STT_TLS` symbol unless there will be
1058 conflicts for that symbol, it doesn't know how many shared libraries with `PT_TLS` segment will precede it or whether
1059 executable will or will not have `PT_TLS` segment. Until TLS is widely deployed by many libraries, prelink could
1060 guess that only `libc.so` will have `PT_TLS` and store 1 (first module ID the dynamic linker assigns), but given that
1061 `libc.so` uses just one such relocation it is not probably worth doing this when soon other shared libraries besides
1062 `libc.so` and `libGL.so` start using it heavily. Because of this prelink doesn't do anything special when prelinking
1063 shared libraries with these relocations and for each relocations in this class creates one conflict fixup.

1064 In second class are relocations which resolve to `st_value` of some `STT_TLS` symbol. These relocations are used in
1065 Global Dynamic TLS model (in Local Dynamic they are resolved at link time already) and from prelink point of
1066 view they are much more similar to normal relocations than the other two classes. When the `STT_TLS` symbol is looked
1067 up successfully in shared library's natural search scope, prelink just stores its `st_value` into the relocation. The
1068 chances there will be a conflict are even smaller than with normal symbol lookups, since overloading TLS symbols
1069 means wasted memory in each single thread and thus library writers will try to avoid it if possible.

1070 The third class includes relocations which resolve to offsets within program's initial TLS block ¹⁵ Relocation in this
1071 class are used in Initial Exec TLS model (or in Local Exec model if this model is supported in shared libraries). These
1072 offsets are even harder to predict than module IDs and unlike module IDs it wouldn't be very helpful if they were
1073 assigned by prelink instead of dynamic linker (which would just read them from some dynamic tag). That's because
1074 TLS block needs to be packed tightly and any assignments in prelink couldn't take into account other shared libraries
1075 linked into the same executable and the executable itself. Similarly to module ID relocations, prelink doesn't do
1076 anything about them when prelinking shared libraries and for each such relocation creates a conflict fixup.

11 Prelinking of executables and shared libraries

1077 Rewriting of executables is harder than for shared libraries, both because there are more changes necessary and because
1078 shared libraries are relocatable and thus have dynamic relocations for all absolute addresses.

1079 After collecting all information from the dynamic linker and assigning virtual address space slots to all shared libraries,
1080 prelinking of shared libraries involves following steps:

- 1081 • Relocation of the shared library to the assigned base address.
- 1082 • REL to RELA conversion if needed (the only step which changes sizes of allocated sections in the middle).
- 1083 • On architectures which have `SHT_NOBITS` `.plt` sections, before relocations are applied the section needs to
1084 be converted to `SHT_PROGBITS`. As the section needs to be at the end (or after it) of file backed part of some
1085 `PT_LOAD` segment, this just means that the file backed up part needs to be enlarged, the file filled with zeros
1086 and all following section file offsets or program header entry file offsets adjusted. All `SHT_NOBITS` sections in
1087 the same `PT_LOAD` segment with virtual addresses lower than the `.plt` start address need to be converted from
1088 `SHT_NOBITS` to `SHT_PROGBITS` too. Without making the section `SHT_PROGBITS`, prelink cannot apply relo-
1089 cations against it as such sections contain only zeros. Architectures with `SHT_NOBITS` `.plt` section supported
1090 by prelink are PowerPC and PowerPC64.

¹⁵Negative on architectures which have TLS block immediately below thread pointer (e.g. IA-32, AMD64, SPARC, S/390) and positive on
architectures which have TLS block at thread pointer or a few bytes above it (e.g. PowerPC, Alpha, IA-64, SuperH).

- Applying relocations. For each dynamic relocation in the shared library, address of relocation's symbol looked up in natural symbol lookup search scope of the shared library (or 0 if the symbol is not found in that search scope) is stored in an architecture and relocation type dependent way to memory pointed by `r_offset` field of the relocation. This step uses symbol lookup information provided by dynamic linker.
- Addition or modification of `DT_CHECKSUM` and `DT_GNU_PRELINKED` dynamic tags.¹⁶ The former is set to checksum of allocated sections in the shared library, the latter to time of prelinking.
- On architectures which don't use writable `.plt`, but instead use `.got=plt` (this section is merged during linking into `.got` section), `prelink` typically stores address into the first PLT slot in `.plt` section to the reserved second word of `.got` section. On these architectures, the dynamic linker has to initialize `.plt` section if lazy binding. On non-prelinked executables or shared libraries this typically means adding load offset to the values in `.got=plt` section, for prelinked shared libraries or executables if prelinking information cannot be used it needs to compute the right values in `.got=plt` section without looking at this section's content (since it contains prelinking information). The second word in `.got` section is used for this computation.
- Addition of `.gnu.prelink.undo` unallocated section if not present yet. This section is used by `prelink` internally during undo operation.
- Addition of `.gnu.liblist` and `.gnu.libstr` unallocated sections or, if they are already present, their update including possible growing or shrinking. These sections are used only by `prelink` to compare the dependent libraries (and their order) at the time when the shared library was prelinked against current dependencies. If a shared library has no dependencies (e.g. dynamic linker), these sections are not present.

Adding or resizing unallocated section needs just file offsets of following unallocated sections recomputed (ensuring proper alignment), growing section header table and `.shstrtab` and adding new section names to that section.

Prelinking of executables involves following steps:

- REL to RELA conversion if needed.
- SHT_NOBITS to SHT_PROGBITS conversion of `.plt` section if needed.
- Applying relocations.
- Addition or resizing of allocated `.gnu.conflict` section containing list of conflict fixups.
- Addition or resizing of allocated `.gnu.liblist` section which is used by the dynamic linker at runtime to see if none of the dependencies changed or were reordered. If they were, it continues normal relocation processing, otherwise they can be skipped and only conflict fixups applied.
- Growing of allocated `.dynstr` section, where strings referenced from `.gnu.liblist` section need to be added.
- If there are any COPY relocations (which `prelink` wants to handle rather than deferring them as conflict fixups to runtime), they need to be applied.
- Modifying second word in `.got` section for `.got=plt` using architectures.
- Addition or adjusting of dynamic tags which allow the dynamic linker to find the `.gnu.liblist` and `.gnu.conflict` sections and their sizes. `DT_GNU_CONFLICT` and `DT_GNU_CONFLICTSZ` should be present if there are any conflict fixups. It should contain the virtual address of the `.gnu.conflict` section start resp. its size in bytes. `DT_GNU_LIBLIST` and `DT_GNU_LIBLISTSZ` need to be present in all prelinked executables and must be equal the to virtual address of the `.gnu.liblist` section and its size in bytes.
- Addition of `.gnu.prelink.undo` unallocated section if not present.

Executables can have absolute relocations already applied (and without a dynamic relocation) to virtually any allocated SHT_PROGBITS section¹⁷, against almost all allocated SHT_PROGBITS and SHT_NOBITS sections. This means that when growing, adding or shrinking allocated sections in executables, all SHT_PROGBITS and SHT_NOBITS section

¹⁶Prelink is not able to grow `.dynamic` section, so it needs some spare dynamic tags (`DT_NULL`) at the end of `.dynamic` section. GNU linker versions released after August 2001 leave space by default.

¹⁷One exception is `.interp` special section. It shouldn't have relocations applied to it, nor any other section should reference it.

1133 must keep their original virtual addresses and sizes¹⁸. Prelink tries various places where to put allocated sections
1134 which were added or grew:

- 1135 • In the unlikely case if there is already some gap between sections in read-only PT_LOAD segment where the
1136 section fits.
- 1137 • If the SHT_NOBITS sections are small enough to fit into a page together with the preceding SHT_PROGBITS
1138 section and there is still some space in the page after the SHT_NOBITS sections. In this case, prelink converts
1139 the SHT_NOBITS sections into SHT_PROGBITS sections, fills them with zeros and adds the new section after it.
1140 This doesn't increase number of PT_LOAD segments, but unfortunately those added sections are writable. This
1141 doesn't matter much for e.g. .gnu.conflict section which is only used before control is transferred to the
1142 program, but could matter for .dynstr which is used even during dlopen.
- 1143 • On IA-32, executables have for historical reasons base address 0x8048000. The reason for this was that when
1144 stack was put immediately below executables, stack and the executable could coexist in the same second level
1145 page table. Linux puts the stack typically at the end of virtual address space and so keeping this exact base
1146 address is not really necessary. Prelink can decrease the base address and thus increase size of read-only
1147 PT_LOAD segment while SHT_PROGBITS and SHT_NOBITS section can stay at their previous addresses. Just their
1148 file offsets need to be increased. All these segment header adjustments need to be done in multiples of ELF
1149 page sizes, so even if prelink chose to do similar things on architectures other than IA-32 which typically
1150 start executables on some address which is a power of 2, it would be only reasonable if ELF page size on that
1151 architecture (which can be much bigger than page size used by the operating system) is very small.
- 1152 • Last possibility is to create a new PT_LOAD segment.¹⁹ Section immediately above program header table
1153 (typically .interp) has to be moved somewhere else, but if possible close to the beginning of the executable.
1154 The new PT_LOAD segment is then added after the last PT_LOAD segment. The segment has to be writable even
1155 when all the sections in it are read-only, unless it ends exactly on a page boundary, because brk area starts
1156 immediately after the end of last PT_LOAD segment and the executable expects it to be writable.

1157 So that verification works properly, if there is .gnu.prelink_undo section in the executable, prelink first reshuffles
1158 the sections and segments for the purpose of finding places for the sections to the original sequence as recorded in the
1159 .gnu.prelink_undo section. Examples of the above mentioned cases:

```
1160 $ SEDCMD='s/^.* \.plt.*$/...;/\[\.\*\.\text\//,\[\.\*\.\got\]/d'  
1161 $ SEDCMD2='/Section to Segment/,\$d;/^Key to/,/^Program/d;/^[\A-Z]/d;/^ *$/d'  
1162 $ cat > test1.c <<EOF  
1163 int main (void) { return 0; }  
1164 EOF  
1165 $ gcc -Wl,--verbose 2>&1 \  
1166 | sed '/^===/,/^===/!d;/^===/d;s/\.\rel\.\dyn/. += 512; &/' > test1.lds  
1167 $ gcc -s -O2 -o test1 test1.c -Wl,-T,test1.lds  
1168 $ readelf -Sl ./test1 | sed -e "$SEDCMD" -e "$SEDCMD2"  
1169 [Nr] Name Type Addr Off Size ES Flg Lk Inf Al  
1170 [ 0] NULL 00000000 000000 000000 00 0 0 0 0  
1171 [ 1] .interp PROGBITS 08048114 000114 000013 00 A 0 0 1  
1172 [ 2] .note.ABI-tag NOTE 08048128 000128 000020 00 A 0 0 4  
1173 [ 3] .hash HASH 08048148 000148 000024 04 A 4 0 4  
1174 [ 4] .dynsym DYNSYM 0804816c 00016c 000040 10 A 5 1 4  
1175 [ 5] .dynstr STRTAB 080481ac 0001ac 000045 00 A 0 0 1  
1176 [ 6] .gnu.version VERSYM 080481f2 0001f2 000008 02 A 4 0 2  
1177 [ 7] .gnu.version_r VERNEED 080481fc 0001fc 000020 00 A 5 1 4  
1178 [ 8] .rel.dyn REL 0804841c 00041c 000008 08 A 4 0 4  
1179 [ 9] .rel.plt REL 08048424 000424 000008 08 A 4 b 4  
1180 [10] .init PROGBITS 0804842c 00042c 000017 00 AX 0 0 4  
1181 ...  
1182 [22] .bss NOBITS 080496f8 0006f8 000004 00 WA 0 0 4
```

¹⁸With a notable exception of splitting one section into two covering the same virtual address range.

¹⁹Linux kernels before 2.4.10 loaded executables which had middle PT_LOAD segment with p_memsz bigger than p_filesz incorrectly, so prelink should be only used on systems with 2.4.10 or later kernels.

```

1183 [ 23] .comment           PROGBITS        00000000 0006f8 000132 00      0  0  1
1184 [ 24] .shstrtab          STRTAB         00000000 00082a 0000be 00      0  0  1
1185 Type      Offset    VirtAddr   PhysAddr  FileSiz MemSiz Flg Align
1186 PHDR      0x000034 0x08048034 0x08048034 0x000e0 0x000e0 R E 0x4
1187 INTERP     0x000114 0x08048114 0x08048114 0x00013 0x00013 R 0x1
1188     [Requesting program interpreter: /lib/ld-linux.so.2]
1189 LOAD       0x000000 0x08048000 0x08048000 0x005fc 0x005fc R E 0x1000
1190 LOAD       0x0005fc 0x080495fc 0x080495fc 0x000fc 0x00100 RW 0x1000
1191 DYNAMIC    0x000608 0x08049608 0x08049608 0x000c8 0x000c8 RW 0x4
1192 NOTE       0x000128 0x08048128 0x08048128 0x00020 0x00020 R 0x4
1193 STACK      0x000000 0x00000000 0x00000000 0x00000 0x00000 RW 0x4
1194 $ prelink -N ./test1
1195 $ readelf -Sl ./test1 | sed -e "$SEDCMD" -e "$SEDCMD2"
1196 [Nr] Name             Type            Addr     Off   Size ES Flg Lk Inf Al
1197 [ 0] NULL             PROGBITS        00000000 00000000 00000000 00      0  0  0
1198 [ 1] .interp          PROGBITS        08048114 000114 000013 00      A  0  0  1
1199 [ 2] .note.ABI-tag   NOTE            08048128 000128 000020 00      A  0  0  4
1200 [ 3] .hash            HASH            08048148 000148 000024 04      A  4  0  4
1201 [ 4] .dynsym          DYNSYM          0804816c 00016c 000040 10      A  8  1  4
1202 [ 5] .gnu.liblist    GNU_LIBLIST     080481ac 0001ac 000028 14      A  8  0  4
1203 [ 6] .gnu.version    VERSYM          080481f2 0001f2 000008 02      A  4  0  2
1204 [ 7] .gnu.version_r  VERNEED         080481fc 0001fc 000020 00      A  8  1  4
1205 [ 8] .dynstr          STRTAB          0804821c 00021c 000058 00      A  0  0  1
1206 [ 9] .gnu.conflict   RELA            08048274 000274 0000c0 0c      A  4  0  4
1207 [10] .rel.dyn         REL             0804841c 00041c 000008 08      A  4  0  4
1208 [11] .rel.plt         REL             08048424 000424 000008 08      A  4  d  4
1209 [12] .init            PROGBITS        0804842c 00042c 000017 00      AX  0  0  4
1210 ...
1211 [ 24] .bss             NOBITS          080496f8 0006f8 000004 00      WA  0  0  4
1212 [ 25] .comment         PROGBITS        00000000 0006f8 000132 00      0  0  1
1213 [ 26] .gnu.prelink_undo PROGBITS        00000000 00082c 0004d4 01      0  0  4
1214 [ 27] .shstrtab         STRTAB          00000000 00d000 0000eb 00      0  0  1
1215 Type      Offset    VirtAddr   PhysAddr  FileSiz MemSiz Flg Align
1216 PHDR      0x000034 0x08048034 0x08048034 0x000e0 0x000e0 R E 0x4
1217 INTERP     0x000114 0x08048114 0x08048114 0x00013 0x00013 R 0x1
1218     [Requesting program interpreter: /lib/ld-linux.so.2]
1219 LOAD       0x000000 0x08048000 0x08048000 0x005fc 0x005fc R E 0x1000
1220 LOAD       0x0005fc 0x080495fc 0x080495fc 0x000fc 0x00100 RW 0x1000
1221 DYNAMIC    0x000608 0x08049608 0x08049608 0x000c8 0x000c8 RW 0x4
1222 NOTE       0x000128 0x08048128 0x08048128 0x00020 0x00020 R 0x4
1223 STACK      0x000000 0x00000000 0x00000000 0x00000 0x00000 RW 0x4

```

Listing 15: Reshuffling of an executable with a gap between sections

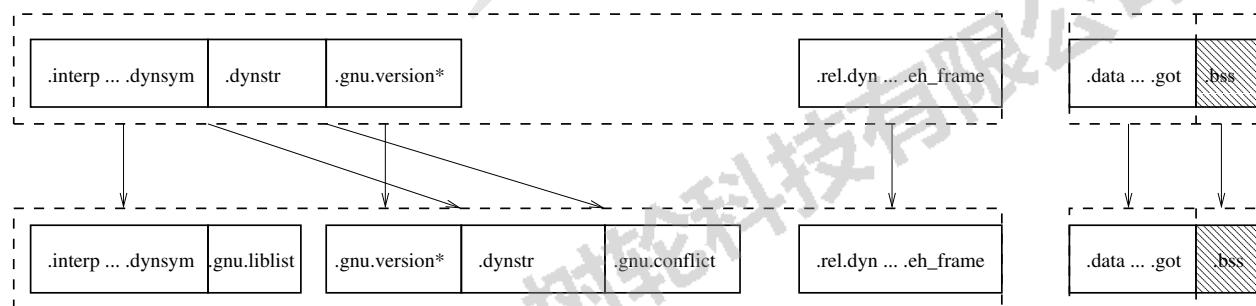


Figure 4: Reshuffling of an executable with a gap between sections

1224 In the above sample, there was enough space between sections (particularly between the end of the .gnu.version_r
1225 section and the start of .rel.dyn) that the new sections could be added there.

1226 \$ SEDCMD='s/^.* \.plt.*\$/.../; /\[.*\.\text/,/\[.*\.\got/d'

```

1227 $ SEDCMD2=' /Section to Segment/, $d; /^Key to/, /^Program/d; /^[A-Z]/d; /^ *$/d'
1228 $ cat > test2.c <<EOF
1229 int main (void) { return 0; }
1230 EOF
1231 $ gcc -s -O2 -o test2 test2.c
1232 $ readelf -Sl ./test2 | sed -e "$SEDCMD" -e "$SEDCMD2"
1233 [Nr] Name           Type      Addr     Off      Size    ES Flg Lk Inf Al
1234 [ 0] .null          NULL      00000000 000000 00000000 00 0 0 0 0
1235 [ 1] .interp        PROGBITS 08048114 000114 000013 00 A 0 0 0 1
1236 [ 2] .note.ABI-tag NOTE      08048128 000128 000020 00 A 0 0 0 4
1237 [ 3] .hash          HASH      08048148 000148 000024 04 A 4 0 0 4
1238 [ 4] .dynsym        DYNAMIC   0804816c 00016c 000040 10 A 5 1 1 4
1239 [ 5] .dynstr        STRTAB   080481ac 0001ac 000045 00 A 0 0 0 1
1240 [ 6] .gnu.version   VERSYM   080481f2 0001f2 000008 02 A 4 0 0 2
1241 [ 7] .gnu.version_r VERNEED   080481fc 0001fc 000020 00 A 5 1 4
1242 [ 8] .rel.dyn       REL      0804821c 00021c 000008 08 A 4 0 0 4
1243 [ 9] .rel.plt       REL      08048224 000224 000008 08 A 4 b 4
1244 [10] .init          PROGBITS 0804822c 00022c 000017 00 AX 0 0 0 4
1245 ...
1246 [22] .bss           NOBITS   080494f8 0004f8 000004 00 WA 0 0 0 4
1247 [23] .comment       PROGBITS 00000000 0004f8 000132 00 0 0 0 1
1248 [24] .shstrtab      STRTAB   00000000 00062a 0000be 00 0 0 0 1
1249 Type      Offset     VirtAddr   PhysAddr   FileSiz MemSiz Flg Align
1250 PHDR     0x000034 0x08048034 0x08048034 0x000e0 0x000e0 R E 0x4
1251 INTERP   0x000114 0x08048114 0x08048114 0x00013 0x00013 R 0x1
1252 [Requesting program interpreter: /lib/ld-linux.so.2]
1253 LOAD     0x000000 0x08048000 0x08048000 0x003fc 0x003fc R E 0x1000
1254 LOAD     0x0003fc 0x080493fc 0x080493fc 0x000fc 0x00100 RW 0x1000
1255 DYNAMIC  0x000408 0x08049408 0x08049408 0x000c8 0x000c8 RW 0x4
1256 NOTE     0x000128 0x08048128 0x08048128 0x00020 0x00020 R 0x4
1257 STACK    0x000000 0x00000000 0x00000000 0x00000 0x00000 RW 0x4
1258 $ prelink -N ./test2
1259 $ readelf -Sl ./test2 | sed -e "$SEDCMD" -e "$SEDCMD2"
1260 [Nr] Name           Type      Addr     Off      Size    ES Flg Lk Inf Al
1261 [ 0] .null          NULL      00000000 000000 00000000 00 0 0 0 0
1262 [ 1] .interp        PROGBITS 08048114 000114 000013 00 A 0 0 0 1
1263 [ 2] .note.ABI-tag NOTE      08048128 000128 000020 00 A 0 0 0 4
1264 [ 3] .hash          HASH      08048148 000148 000024 04 A 4 0 0 4
1265 [ 4] .dynsym        DYNAMIC   0804816c 00016c 000040 10 A 23 1 1 4
1266 [ 5] .gnu.liblist   GNU_LIBLIST 080481ac 0001ac 000028 14 A 23 0 0 4
1267 [ 6] .gnu.version   VERSYM   080481f2 0001f2 000008 02 A 4 0 0 2
1268 [ 7] .gnu.version_r VERNEED   080481fc 0001fc 000020 00 A 23 1 4
1269 [ 8] .rel.dyn       REL      0804821c 00021c 000008 08 A 4 0 0 4
1270 [ 9] .rel.plt       REL      08048224 000224 000008 08 A 4 b 4
1271 [10] .init          PROGBITS 0804822c 00022c 000017 00 AX 0 0 0 4
1272 ...
1273 [22] .bss           PROGBITS 080494f8 0004f8 000004 00 WA 0 0 0 4
1274 [23] .dynstr        STRTAB   080494fc 0004fc 000058 00 A 0 0 0 1
1275 [24] .gnu.conflict  RELA     08049554 000554 0000c0 0c A 4 0 0 4
1276 [25] .comment       PROGBITS 00000000 000614 000132 00 0 0 0 1
1277 [26] .gnu.prelink_undo PROGBITS 00000000 000748 0004d4 01 0 0 0 4
1278 [27] .shstrtab      STRTAB   00000000 000c1c 0000eb 00 0 0 0 1
1279 Type      Offset     VirtAddr   PhysAddr   FileSiz MemSiz Flg Align
1280 PHDR     0x000034 0x08048034 0x08048034 0x000e0 0x000e0 R E 0x4
1281 INTERP   0x000114 0x08048114 0x08048114 0x00013 0x00013 R 0x1
1282 [Requesting program interpreter: /lib/ld-linux.so.2]
1283 LOAD     0x000000 0x08048000 0x08048000 0x003fc 0x003fc R E 0x1000
1284 LOAD     0x0003fc 0x080493fc 0x080493fc 0x00218 0x00218 RW 0x1000
1285 DYNAMIC  0x000408 0x08049408 0x08049408 0x000c8 0x000c8 RW 0x4
1286 NOTE     0x000128 0x08048128 0x08048128 0x00020 0x00020 R 0x4
1287 STACK    0x000000 0x00000000 0x00000000 0x00000 0x00000 RW 0x4

```

Listing 16: Reshuffling of an executable with small .bss

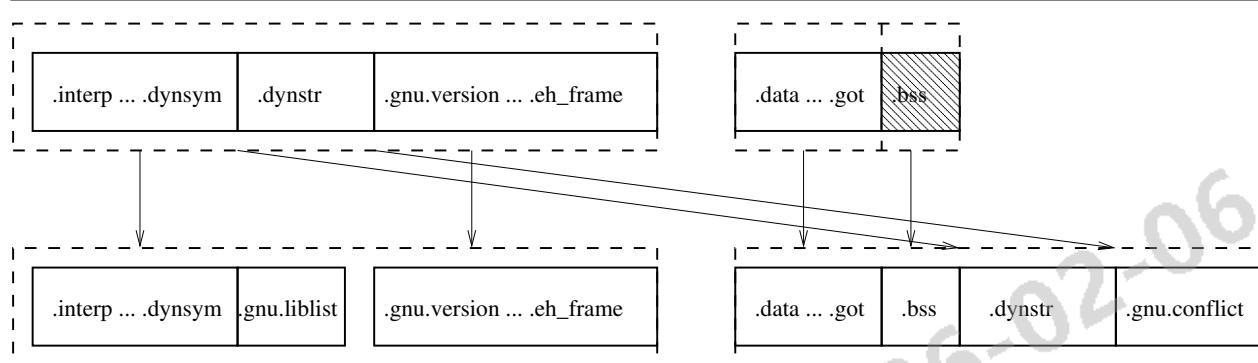


Figure 5: Reshuffling of an executable with small .bss

¹²⁸⁸ In this case .bss section was small enough that prelink converted it to SHT_PROGBITS.

```

1289 $ SEDCMD='s/^.* \.plt.*$/...;/\[^.\].text//;\[^.\].got/d'
1290 $ SEDCMD2='/Section to Segment/,${d};/^Key to/,^Program/d;^[A-Z]/d;^ *$/{d'
1291 $ cat > test3.c <<EOF
1292 int foo [4096];
1293 int main (void) { return 0; }
1294 EOF
1295 $ gcc -s -O2 -o test3 test3.c
1296 $ readelf -Sl ./test3 | sed -e "$SEDCMD" -e "$SEDCMD2"
1297 [Nr] Name           Type        Addr     Off      Size   ES Flg Lk Inf Al
1298 [ 0] NULL           PROGBITS    00000000 000000 000000 00  0  0  0  0
1299 [ 1] .interp        PROGBITS    08048114 000114 000013 00  A  0  0  1
1300 [ 2] .note.ABI-tag NOTE        08048128 000128 000020 00  A  0  0  4
1301 [ 3] .hash          HASH        08048148 000148 000024 04  A  4  0  4
1302 [ 4] .dynsym        DYNSYM     0804816c 00016c 000040 10  A  5  1  4
1303 [ 5] .dynstr        STRTAB     080481ac 0001ac 000045 00  A  0  0  1
1304 [ 6] .gnu.version   VERSYM     080481f2 0001f2 000008 02  A  4  0  2
1305 [ 7] .gnu.version_r VERNEED    080481fc 0001fc 000020 00  A  5  1  4
1306 [ 8] .rel.dyn       REL        0804821c 00021c 000008 08  A  4  0  4
1307 [ 9] .rel.plt      REL        08048224 000224 000008 08  A  4  b  4
1308 [10] .init          PROGBITS   0804822c 00022c 000017 00  AX  0  0  4
1309 ...
1310 [22] .bss           NOBITS     08049500 000500 004020 00  WA  0  0 32
1311 [23] .comment       PROGBITS   00000000 000500 000132 00  0  0  1
1312 [24] .shstrtab      STRTAB     00000000 000632 0000be 00  0  0  1
1313 Type            Offset     VirtAddr   PhysAddr  FileSiz MemSiz Flg Align
1314 PHDR           0x000034 0x08048034 0x08048034 0x000e0 0x000e0 R E 0x4
1315 INTERP          0x000114 0x08048114 0x08048114 0x00013 0x00013 R 0x1
1316 [Requesting program interpreter: /lib/ld-linux.so.2]
1317 LOAD            0x000000 0x08048000 0x08048000 0x003fc 0x003fc R E 0x1000
1318 LOAD            0x0003fc 0x080493fc 0x080493fc 0x000fc 0x04124 RW 0x1000
1319 DYNAMIC         0x000408 0x08049408 0x08049408 0x000c8 0x000c8 RW 0x4
1320 NOTE            0x000128 0x08048128 0x08048128 0x00020 0x00020 R 0x4
1321 STACK           0x000000 0x00000000 0x00000000 0x000000 0x000000 RW 0x4
1322 $ prelink -N ./test3
1323 $ readelf -Sl ./test3 | sed -e "$SEDCMD" -e "$SEDCMD2"
1324 [Nr] Name           Type        Addr     Off      Size   ES Flg Lk Inf Al
1325 [ 0] NULL           PROGBITS   00000000 000000 000000 00  0  0  0
1326 [ 1] .interp        PROGBITS   08047114 000114 000013 00  A  0  0  1
1327 [ 2] .note.ABI-tag NOTE        08047128 000128 000020 00  A  0  0  4
1328 [ 3] .dynstr        STRTAB     08047148 000148 000058 00  A  0  0  1
1329 [ 4] .gnu.liblist   GNU_LIBLIST 080471a0 0001a0 000028 14  A  3  0  4
1330 [ 5] .gnu.conflict RELA        080471c8 0001c8 0000c0 0c  A  7  0  4
1331 [ 6] .hash          HASH        08048148 001148 000024 04  A  7  0  4
1332 [ 7] .dynsym        DYNSYM     0804816c 00116c 000040 10  A  3  1  4
1333 [ 8] .gnu.version   VERSYM     080481f2 0011f2 000008 02  A  7  0  2

```

```

1334 [ 9] .gnu.version_r VERNEED 080481fc 0011fc 000020 00 A 3 1 4
1335 [10] .rel.dyn REL 0804821c 00121c 000008 08 A 7 0 4
1336 [11] .rel.plt REL 08048224 001224 000008 08 A 7 d 4
1337 [12] .init PROGBITS 0804822c 00122c 000017 00 AX 0 0 4
1338 ...
1339 [24] .bss NOBITS 08049500 0014f8 004020 00 WA 0 0 32
1340 [25] .comment PROGBITS 00000000 0014f8 000132 00 0 0 1
1341 [26] .gnu.prelink_undo PROGBITS 00000000 00162c 0004d4 01 0 0 4
1342 [27] .shstrtab STRTAB 00000000 001b00 0000eb 00 0 0 1
1343 Type Offset VirtAddr PhysAddr FileSiz MemSiz Flg Align
1344 PHDR 0x000034 0x08047034 0x08047034 0x000e0 0x000e0 R E 0x4
1345 INTERP 0x000114 0x08047114 0x08047114 0x00013 0x00013 R 0x1
1346 [Requesting program interpreter: /lib/ld-linux.so.2]
1347 LOAD 0x000000 0x08047000 0x08047000 0x013fc 0x013fc R E 0x1000
1348 LOAD 0x0013fc 0x080493fc 0x080493fc 0x000fc 0x04124 RW 0x1000
1349 DYNAMIC 0x001408 0x08049408 0x08049408 0x000c8 0x000c8 RW 0x4
1350 NOTE 0x000128 0x08047128 0x08047128 0x00020 0x00020 R 0x4
1351 STACK 0x000000 0x00000000 0x00000000 0x000000 0x000000 RW 0x4

```

Listing 17: Reshuffling of an executable with decreasing of base address

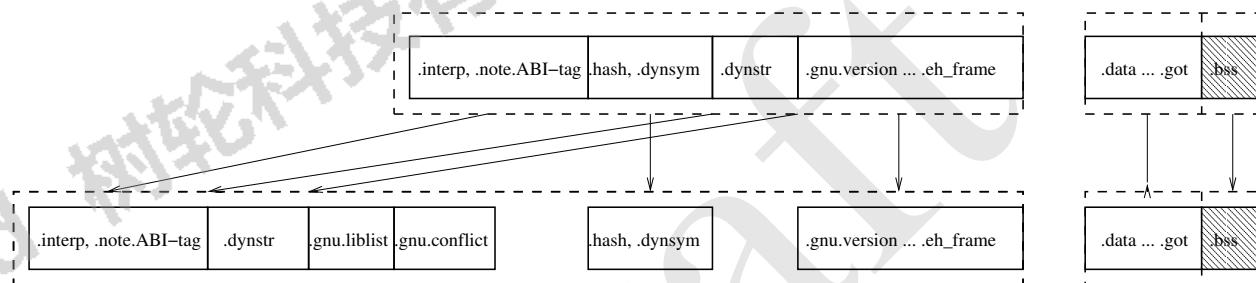


Figure 6: Reshuffling of an executable with decreasing of the base address

1352 In test3 the base address of the executable was decreased by one page and the new sections added there.

```

1353 $ SEDCMD='s/^.* \.plt.*$/.../;/\[.*\.\text//,\[.*\.\got/d'
1354 $ SEDCMD2='/Section to Segment/,\$;/^Key to/,/^Program/d;/^ [A-Z]/d;/^ *$/d'
1355 $ cat > test4.c <<EOF
1356 int foo [4096];
1357 int main (void) { return 0; }
1358 EOF
1359 $ gcc -Wl,--verbose 2>&1 \
1360 | sed '/^===/,/^===/!d;/^==/d;s/0x08048000/0x08000000/' > test4.lds
1361 $ gcc -s -O2 -o test4 test4.c -Wl,-T,test4.lds
1362 $ readelf -Sl ./test4 | sed -e "$SEDCMD" -e "$SEDCMD2"
1363 [Nr] Name Type Addr Off Size ES Flg Lk Inf Al
1364 [ 0] NULL NULL 00000000 000000 000000 00 0 0 0
1365 [ 1] .interp PROGBITS 08000114 000114 000013 00 A 0 0 1
1366 [ 2] .note.ABI-tag NOTE 08000128 000128 000020 00 A 0 0 4
1367 [ 3] .hash HASH 08000148 000148 000024 04 A 4 0 4
1368 [ 4] .dynsym DYNSYM 0800016c 00016c 000040 10 A 5 1 4
1369 [ 5] .dynstr STRTAB 080001ac 0001ac 000045 00 A 0 0 1
1370 [ 6] .gnu.version VERSYM 080001f2 0001f2 000008 02 A 4 0 2
1371 [ 7] .gnu.version_r VERNEED 080001fc 0001fc 000020 00 A 5 1 4
1372 [ 8] .rel.dyn REL 0800021c 00021c 000008 08 A 4 0 4
1373 [ 9] .rel.plt REL 08000224 000224 000008 08 A 4 b 4
1374 [10] .init PROGBITS 0800022c 00022c 000017 00 AX 0 0 4
1375 ...
1376 [22] .bss NOBITS 08001500 000500 004020 00 WA 0 0 32
1377 [23] .comment PROGBITS 00000000 000500 000132 00 0 0 1

```

```

1378 [ 24] .shstrtab      STRTAB      00000000 000632 0000be 00      0  0  1
1379 Type          Offset  VirtAddr  PhysAddr  FileSiz MemSiz Flg Align
1380 PHDR          0x000034 0x80000034 0x80000034 0x000e0 0x000e0 R E 0x4
1381 INTERP         0x000114 0x80000114 0x80000114 0x00013 0x00013 R  0x1
1382 [Requesting program interpreter: /lib/ld-linux.so.2]
1383 LOAD           0x000000 0x80000000 0x80000000 0x003fc 0x003fc R E 0x1000
1384 LOAD           0x0003fc 0x80013fc 0x80013fc 0x000fc 0x04124 RW  0x1000
1385 DYNAMIC        0x000408 0x8001408 0x8001408 0x000c8 0x000c8 RW  0x4
1386 NOTE           0x000128 0x80000128 0x80000128 0x00020 0x00020 R  0x4
1387 STACK          0x000000 0x00000000 0x00000000 0x00000 0x00000 RW  0x4
1388 $ prelink -N ./test4
1389 $ readelf -Sl ./test4 | sed -e "$SEDCMD" -e "$SEDCMD2"
1390 [Nr] Name          Type          Addr      Off      Size   ES Flg Lk Inf Al
1391 [ 0] NULL          NULL          00000000 000000 00000000 00      0  0  0
1392 [ 1] .interp       PROGBITS    08000134 000134 000013 00      A  0  0  1
1393 [ 2] .note.ABI-tag NOTE          08000148 000148 000020 00      A  0  0  4
1394 [ 3] .hash          HASH          08000168 000168 000024 04      A  4  0  4
1395 [ 4] .dynsym        DYNSYM        0800018c 00018c 000040 10      A 22  1  4
1396 [ 5] .gnu.version   VERSYM        080001f2 0001f2 000008 02      A  4  0  2
1397 [ 6] .gnu.version_r VERNEED      080001fc 0001fc 000020 00      A 22  1  4
1398 [ 7] .rel.dyn       REL           0800021c 00021c 000008 08      A  4  0  4
1399 [ 8] .rel.plt      REL           08000224 000224 000008 08      A  4  a  4
1400 [ 9] .init          PROGBITS    0800022c 00022c 000017 00      AX  0  0  4
1401 ...
1402 [ 21] .bss          NOBITS       08001500 0004f8 004020 00      WA  0  0 32
1403 [ 22] .dynstr       STRTAB       080064f8 0004f8 000058 00      A  0  0  1
1404 [ 23] .gnu.liblist   GNU_LIBLIST 08006550 000550 000028 14      A 22  0  4
1405 [ 24] .gnu.conflict  RELA          08006578 000578 0000c0 0c      A  4  0  4
1406 [ 25] .comment      PROGBITS    00000000 000638 000132 00      0  0  1
1407 [ 26] .gnu.prelink_undo PROGBITS  00000000 00076c 0004d4 01      0  0  4
1408 [ 27] .shstrtab      STRTAB       00000000 000c40 0000eb 00      0  0  1
1409 Type          Offset  VirtAddr  PhysAddr  FileSiz MemSiz Flg Align
1410 PHDR          0x000034 0x80000034 0x80000034 0x000e0 0x000e0 R E 0x4
1411 INTERP         0x000134 0x80000134 0x80000134 0x00013 0x00013 R  0x1
1412 [Requesting program interpreter: /lib/ld-linux.so.2]
1413 LOAD           0x000000 0x80000000 0x80000000 0x003fc 0x003fc R E 0x1000
1414 LOAD           0x0003fc 0x80013fc 0x80013fc 0x000fc 0x04124 RW  0x1000
1415 LOAD           0x0004f8 0x80064f8 0x80064f8 0x00140 0x00140 RW  0x1000
1416 DYNAMIC        0x000408 0x8001408 0x8001408 0x000c8 0x000c8 RW  0x4
1417 NOTE           0x000148 0x80000148 0x80000148 0x00020 0x00020 R  0x4
1418 STACK          0x000000 0x00000000 0x00000000 0x00000 0x00000 RW  0x4

```

Listing 18: Reshuffling of an executable with addition of a new segment

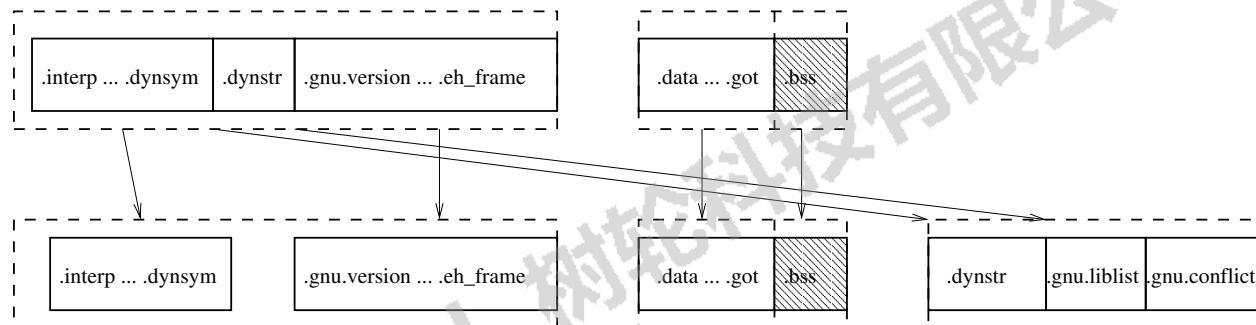


Figure 7: Reshuffling of an executable with addition of a new segment

¹⁴¹⁹ In the last example, base address was not decreased but instead a new PT_LOAD segment has been added.

¹⁴²⁰ R_<arch>.COPY relocations are typically against first part of the SHT_NOBITS .bss section. So that prelink can ¹⁴²¹ apply them, it needs to first change their section to SHT_PROGBITS, but as .bss section typically occupies much larger

part of memory, it is not desirable to convert .bss section into SHT_PROGBITS as whole. A section cannot be partly SHT_PROGBITS and partly SHT_NOBITS, so prelink first splits the section into two parts, first .dynbss which covers area from the start of .bss section up to highest byte to which some COPY relocation is applied and then the old .bss. The first is converted to SHT_PROGBITS and its size is decreased, the latter stays SHT_NOBITS and its start address and file offset are adjusted as well as its size decreased. The dynamic linker handles relocations in the executable last, so prelink cannot just copy memory from the shared library where the symbol of the COPY relocation has been looked up in. There might be relocations applied by the dynamic linker in normal relocation processing to the objects, so prelink has to first process the relocations against that memory area. Relocations which don't need conflict fixups are already applied, so prelink just needs to apply conflict fixups against the memory area, then copy it to the newly created .dynbss section.

Here is an example which shows various things which COPY relocation handling in prelink needs to deal with:

```
1433 $ cat > test1.c <<EOF
1434 struct A { char a; struct A *b; int *c; int *d; };
1435 int bar, baz;
1436 struct A foo = { 1, &foo, &bar, &baz };
1437 int *addr (void) { return &baz; }
1438 EOF
1439 $ cat > test.c <<EOF
1440 #include <stdio.h>
1441 struct A { char a; struct A *b; int *c; int *d; };
1442 int bar, *addr (void), big[8192];
1443 extern struct A foo;
1444 int main (void)
1445 {
1446     printf ("%p: %d %p %p %p %p\n", &foo, foo.a, foo.b, foo.c, foo.d,
1447     &bar, addr ());
1448 }
1449 EOF
1450 $ gcc -nostdlib -shared -fpic -s -o test1.so test1.c
1451 $ gcc -s -o test test.c ./test1.so
1452 $ ./test
1453 0x80496c0: 1 0x80496c0 0x80516e0 0x4833a4 0x80516e0 0x4833a4
1454 $ readelf -r test | sed '/\.rel\.\dyn/,/\.rel\.\plt/!d;/^0/!d'
1455 080496ac 00000c06 R_386_GLOB_DAT 00000000 __gmon_start__
1456 080496c0 00000605 R_386_COPY 080496c0 foo
1457 $ readelf -S test | grep bss
1458 [22] .bss NOBITS 080496c0 0006c0 008024 00 WA 0 0 32
1459 $ prelink -N ./test ./test1.so
1460 $ readelf -s test | grep foo
1461 6: 080496c0 16 OBJECT GLOBAL DEFAULT 25 foo
1462 $ readelf -s test1.so | grep foo
1463 15: 004a9314 16 OBJECT GLOBAL DEFAULT 6 foo
1464 $ readelf -r test | sed '/.gnu.conflict/,/\.rel\.\dyn/!d;/^0/!d'
1465 004a9318 00000001 R_386_32 080496c0
1466 004a931c 00000001 R_386_32 080516e0
1467 005f9874 00000001 R_386_32 fffffff0
1468 005f9878 00000001 R_386_32 00000001
1469 005f98bc 00000001 R_386_32 fffffff4
1470 005f9900 00000001 R_386_32 ffffffec
1471 005f9948 00000001 R_386_32 ffffffdc
1472 005f995c 00000001 R_386_32 ffffffe0
1473 005f9980 00000001 R_386_32 fffffff8
1474 005f9988 00000001 R_386_32 ffffffe4
1475 005f99a4 00000001 R_386_32 ffffffd8
1476 005f99c4 00000001 R_386_32 ffffffe8
1477 005f99d8 00000001 R_386_32 08048584
1478 004c2510 00000007 R_386_JUMP_SLOT 00534460
1479 004c2514 00000007 R_386_JUMP_SLOT 00534080
1480 004c2518 00000007 R_386_JUMP_SLOT 00534750
1481 004c251c 00000007 R_386_JUMP_SLOT 005342c0
```

```

1482 004c2520 00000007 R_386_JUMP_SLOT 00534200
1483 $ objdump -s -j .dynbss test
1484
1485 test:      file format elf32-i386
1486
1487 Contents of section .dynbss:
1488 80496c0 01000000 c0960408 e0160508 a4934a00 .....J.
1489 $ objdump -s -j .data test1.so
1490
1491 test1.so:    file format elf32-i386
1492
1493 Contents of section .data:
1494 4a9314 01000000 14934a00 a8934a00 a4934a00 .....J....J....J.
1495 $ readelf -S test | grep bss
1496 [ 24] .dynbss      PROGBITS 080496c0 0016c0 000010 00 WA 0 0 32
1497 [ 25] .bss        NOBITS 080496d0 0016d0 008014 00 WA 0 0 32
1498 $ sed 's/8192/1/' test.c > test2.c
1499 $ gcc -s -o test2 test2.c ./test1.so
1500 $ readelf -S test2 | grep bss
1501 [ 22] .bss        NOBITS 080496b0 0006b0 00001c 00 WA 0 0 8
1502 $ prelink -N ./test2 ./test1.so
1503 $ readelf -S test2 | grep bss
1504 [ 22] .dynbss      PROGBITS 080496b0 0006b0 000010 00 WA 0 0 8
1505 [ 23] .bss        PROGBITS 080496c0 0006c0 00000c 00 WA 0 0 8

```

Listing 19: Relocation handling of .dynbss objects

1506 Because `test.c` executable is not compiled as position independent code and takes address of `foo` variable, a COPY
 1507 relocation is needed to avoid dynamic relocation against executable's read-only PT_LOAD segment. The `foo` object
 1508 in `test1.so` has one field with no relocations applied at all, one relocation against the variable itself, one relocation
 1509 which needs a conflict fixup (as it is overridden by the variable in the executable) and one with relocation which doesn't
 1510 need any fixups. The first and last field contain already the right values in prelinked `test1.so`, while second and third
 1511 one need to be changed for symbol addresses in the executable (as shown in the `objdump` output). The conflict fixups
 1512 against `foo` in `test1.so` need to stay (unless it is a C++ virtual table or RTTI data, i.e. not in this testcase). In
 1513 `test`, `prelink` changed `.dynbss` to SHT_PROGBITS and kept SHT_NOBITS `.bss`, while in slightly modified testcase
 1514 (`test2`) the size of `.bss` was small enough that `prelink` chose to make it SHT_PROGBITS too and grow the read-write
 1515 PT_LOAD segment and put `.dynstr` and `.gnu.conflict` sections after it.

12 Prelink undo operation

1516 Prelinking of shared libraries and executables is designed to be reversible, so that prelink operation followed by undo
 1517 operation generates bitwise identical file to the original before prelinking. For this operation `prelink` stores the orig-
 1518 inal ELF header, all the program and all section headers into a `.gnu.prelink.undo` section before it starts prelinking
 1519 an unprelinked executable or shared library. When undoing the modifications, `prelink` has to convert RELA back
 1520 to REL first if REL to RELA conversion was done during prelinking and all allocated sections above it relocated down
 1521 to adjust for the section shrink. Relocation types which were changed when trying to avoid REL to RELA conversion
 1522 need to be changed back (e.g. on IA-32, it is assumed R_386_GLOB_DAT relocations should be only those against `.got`
 1523 section and R_386_32 relocations in the remaining places). On RELA architectures, the memory pointed by `r_offset`
 1524 field of the relocations needs to be reinitialized to the values stored there by the linker originally. For `prelink` it
 1525 doesn't matter much what this value is (e.g. always 0, copy of `r_addend`, etc.), as long as it is computable from the
 1526 information `prelink` has during undo operation²⁰. The GNU linker had to be changed on several architectures, so
 1527 that it stores there such a value, as in several places the value e.g. depended on original addend before final link (which
 1528 is not available anywhere after final link time, since `r_addend` field could be adjusted during the final link). If second
 1529 word of `.got` section has been modified, it needs to be reverted back to the original value (on most architectures zero).
 1530 In executables, sections which were moved during prelinking need to be put back and segments added while prelinking
 1531 must be removed.

²⁰Such as relocation type, `r_addend` value, type, binding, flags or other attributes of relocation's symbol, what section the relocation points into or the offset within section it points to.

1532 There are 3 different ways how an undo operation can be performed:

- 1533 • Undoing individual executables or shared libraries specified on the command line in place (i.e. when the undo
1534 operation is successful, the prelinked executable or library is atomically replaced with the undone object).
- 1535 • With `-o` option, only a single executable or shared library given on the command line is undone and stored to the
1536 file specified as `-o` option's argument.
- 1537 • With `-ua` options, `prelink` builds a list of executables in paths written in its config file (plus directories and
1538 executables or libraries from command line) and all shared libraries these executables depend on. All executables
1539 and libraries in the list are then unprelinked. This option is used to unprelink the whole system. It is not perfect
1540 and needs to be worked on, since e.g. if some executable uses some shared library which no other executable
1541 links against, this executable (and shared library) is prelinked, then the executable is removed (e.g. uninstalled)
1542 but the shared library is kept, then the shared library is not unprelinked unless specifically mentioned on the
1543 command line.

13 Verification of prelinked files

1544 As `prelink` needs to modify executables and shared libraries installed on a system, it complicates system integrity
1545 verification (e.g. `rpm -V`, TripWire). These systems store checksums of installed files into some database and during
1546 verification compute them again and compare to the values stored in the database. On a prelinked system most of the
1547 executables and shared libraries would be reported as modified. `Prelink` offers a special mode for these systems, in
1548 which it verifies that unprelinking the executable or shared library followed by immediate prelinking (with the same
1549 base address) creates bitwise identical output with the executable or shared library that's being verified. Furthermore,
1550 depending on other `prelink` options, it either writes the unprelinked image to its standard output or computes MD5
1551 or SHA1 digest from this unprelinked image. Mere undo operation to a file and checksumming it is not good enough,
1552 since an intruder could have modified e.g. conflict fixups or memory which relocations point at, changing a behavior
1553 of the program while file after unprelinking would be unmodified.

1554 During verification, both `prelink` executable and the dynamic linker are used, so a proper system integrity verifica-
1555 tion first checks whether `prelink` executable (which is statically linked for this reason) hasn't been modified, then
1556 uses `prelink --verify` to verify the dynamic linker (when verifying `ld.so` the dynamic linker is not executed)
1557 followed by verification of other executables and libraries.

1558 Verification requires all dependencies of checked object to be unmodified since last prelinking. If some dependency
1559 has been changed or is missing, `prelink` will report it and return with non-zero exit status. This is because prelinking
1560 depends on their content and so if they are modified, the executable or shared library might be different to one after
1561 unprelinking followed by prelinking again. In the future, perhaps it would be possible to even verify executables or
1562 shared libraries without unmodified dependencies, under the assumption that in such case the prelink information will
1563 not be used. It would just need to verify that nothing else but the information only used when dependencies are up to
1564 date has changed between the executable or library on the filesystem and file after unprelink followed by prelink cycle.
1565 The `prelink` operation would need to be modified in this case, so that no information is collected from the dynamic
1566 linker, the list of dependencies is assumed to be the one stored in the executable and expect it to have identical number
1567 of conflict fixups.

14 Measurements

1568 There are two areas where `prelink` can speed things up noticeably. The primary is certainly startup time of big GUI
1569 applications where the dynamic linker spends from 100ms up to a few seconds before giving control to the application.
1570 Another area is when lots of small programs are started up, but their execution time is rather short, so the startup time
1571 which `prelink` optimizes is a noticeable fraction of the total time. This is typical for shell scripting.

1572 First numbers are from `lmbench` benchmark, version 3.0-a3. Most of the benchmarks in `lmbench` suite measure kernel
1573 speed, so it doesn't matter much whether `prelink` is used or not. Only in `lat_proc` benchmark `prelink` shows up
1574 visibly. This benchmark measures 3 different things:

- 1575 • `fork proc`, which is `fork()` followed by immediate `exit(1)` in the child and `wait(0)` in the parent. The results
1576 are (as expected) about the same between unprelinked and prelinked systems.

-
- 1577 • exec proc, i.e. fork() followed by immediate close(1) and execve() of a simple hello world program
 1578 (this program is compiled and linked during the benchmark into a temporary directory and is never prelinked).
 1579 The numbers are 160 μ s to 200 μ s better on prelinked systems, because there is no relocation processing needed
 1580 initially in the dynamic linker and because all relative relocations in libc.so.6 can be skipped.
- 1581 • sh proc, i.e. fork() followed by immediate close(1) and execlp("/bin/sh", "sh", "-c", "/tmp/hello",
 1582 0). Although the hello world program is not prelinked in this case either, the shell is, so out of the 900 μ s to
 1583 1000 μ s speedup less than 200 μ s can be accounted on the speed up of the hello world program as in exec proc
 1584 benchmark and the rest to the speedup of shell startup.

1585 First 4 rows are from running the benchmark on a fully unprelinked system, the last 4 rows on the same system, but
 1586 fully prelinked.

```
1587          L M B E N C H  3 . 0   S U M M A R Y
1588  -----
1589  (Alpha software, do not distribute)
1590
1591 Processor, Processes - times in microseconds - smaller is better
1592  -----
1593 Host          OS  Mhz null null      open slct sig  sig  fork exec sh
1594           call  I/O stat clos TCP  inst hndl proc proc proc
1595  -----
1596 pork Linux 2.4.22 651 0.53 0.97 6.20 8.10 41.2 1.44 4.30 276. 1497 5403
1597 pork Linux 2.4.22 651 0.53 0.95 6.14 7.91 37.8 1.43 4.34 274. 1486 5391
1598 pork Linux 2.4.22 651 0.56 0.94 6.18 8.09 43.4 1.41 4.30 251. 1507 5423
1599 pork Linux 2.4.22 651 0.53 0.94 6.12 8.09 41.0 1.43 4.40 256. 1497 5385
1600 pork Linux 2.4.22 651 0.56 0.94 5.79 7.58 39.1 1.41 4.30 271. 1319 4460
1601 pork Linux 2.4.22 651 0.56 0.92 5.76 7.40 38.9 1.41 4.30 253. 1304 4417
1602 pork Linux 2.4.22 651 0.56 0.95 6.20 7.83 37.7 1.41 4.37 248. 1323 4481
1603 pork Linux 2.4.22 651 0.56 1.01 6.04 7.77 37.9 1.43 4.32 256. 1324 4457
```

Listing 20: lmbench results without and with prelinking

1604 Below is a sample timing of a 239K long configure shell script from GCC on both unprelinked and prelinked system.
 1605 Preparation step was following:

```
1606 cd; cvs -d :pserver:anoncvs@subversions.gnu.org:/cvsroot/gcc login
1607 # Empty password
1608 cvs -d :pserver:anoncvs@subversions.gnu.org:/cvsroot/gcc -z3 co -D20031103 gcc
1609 mkdir ~./gcc/obj
1610 cd ~./gcc/obj; ./configure i386-redhat-linux; make configure-gcc
```

Listing 21: Preparation script for shell script tests

1611 On an unprelinked system, the results were:

```
1612 cd ~./gcc/obj/gcc
1613 for i in 1 2; do ./config.status --recheck > /dev/null 2>&1; done
1614 for i in 1 2 3 4; do time ./config.status --recheck > /dev/null 2>&1; done
1615
1616 real    0m4.436s
1617 user    0m1.730s
1618 sys     0m1.260s
1619
```

```
1620 real    0m4.409s
1621 user    0m1.660s
1622 sys     0m1.340s
1623
1624 real    0m4.431s
1625 user    0m1.810s
1626 sys     0m1.300s
1627
1628 real    0m4.432s
1629 user    0m1.670s
1630 sys     0m1.210s
```

Listing 22: Shell script test results on unprelinked system

1631 and on a fully prelinked system:

```
1632 cd ~/gcc/obj/gcc
1633 for i in 1 2; do ./config.status --recheck > /dev/null 2>&1; done
1634 for i in 1 2 3 4; do time ./config.status --recheck > /dev/null 2>&1; done
1635
1636 real    0m4.126s
1637 user    0m1.590s
1638 sys     0m1.240s
1639
1640 real    0m4.151s
1641 user    0m1.620s
1642 sys     0m1.230s
1643
1644 real    0m4.161s
1645 user    0m1.600s
1646 sys     0m1.190s
1647
1648 real    0m4.122s
1649 user    0m1.570s
1650 sys     0m1.230s
```

Listing 23: Shell script test results on prelinked system

1651 Now timing of a few big GUI programs. All timings were done without X server running and with DISPLAY environment variable not set (so that when control is transferred to the application, it very soon finds out there is no X server 1652 it can talk to and bail out). The measurements are done by the dynamic linker in ticks on a 651MHz dual Pentium III 1653 machine, i.e. ticks have to be divided by 651000000 to get times in seconds. Each application has been run 4 times and 1654 the results with smallest total time spent in the dynamic linker was chosen. Epiphany WWW browser and Evolution 1655 mail client were chosen as examples of G_{tk}+ applications (typically they use really many shared libraries, but many 1656 of them are quite small, there aren't really many relocations nor conflict fixups and most of the libraries are written 1657 in C) and Konqueror WWW browser and KWord word processor were chosen as examples of KDE applications (typ- 1658 ically they use slightly fewer shared libraries, though still a lot, most of the shared libraries are written in C++, have 1659 many relocations and cause many conflict fixups, especially without C++ conflict fixup optimizations in prelink). 1660 On non-prelinked system, the timings are done with lazy binding, i.e. without LD_BIND_NOW=1 set in the environment. 1661 This is because that's how people generally run programs, on the other side it is not exact apples to apples comparison, 1662 since on prelinked system there is no lazy binding with the exception of shared libraries loaded through dlopen. So 1663 when control is passed to the application, prelinked programs should be slightly faster for a while since non-prelinked 1664 programs will have to do symbol lookups and processing relocations (and on various architectures flushing instruction 1665 caches) whenever they call some function they haven't called before in particular shared library or in the executable.

```
1667 $ ldd `which epiphany-bin` | wc -l
```

```
1668      64
1669 $ # Unprelinked system
1670 $ LD_DEBUG=statistics epiphany-bin 2>&1 | sed 's/^ *//'
1671 18960:
1672 18960:    runtime linker statistics:
1673 18960:        total startup time in dynamic loader: 67336593 clock cycles
1674 18960:            time needed for relocation: 58119983 clock cycles (86.3%)
1675 18960:            number of relocations: 6999
1676 18960:            number of relocations from cache: 4770
1677 18960:            number of relative relocations: 31494
1678 18960:            time needed to load objects: 8696104 clock cycles (12.9%)
1679
1680 (epiphany-bin:18960): Gtk-WARNING **: cannot open display:
1681 18960:
1682 18960:    runtime linker statistics:
1683 18960:        final number of relocations: 7692
1684 18960:        final number of relocations from cache: 4770
1685 $ # Prelinked system
1686 $ LD_DEBUG=statistics epiphany-bin 2>&1 | sed 's/^ *//'
1687 25697:
1688 25697:    runtime linker statistics:
1689 25697:        total startup time in dynamic loader: 7313721 clock cycles
1690 25697:            time needed for relocation: 565680 clock cycles (7.7%)
1691 25697:            number of relocations: 0
1692 25697:            number of relocations from cache: 1205
1693 25697:            number of relative relocations: 0
1694 25697:            time needed to load objects: 6179467 clock cycles (84.4%)
1695
1696 (epiphany-bin:25697): Gtk-WARNING **: cannot open display:
1697 25697:
1698 25697:    runtime linker statistics:
1699 25697:        final number of relocations: 31
1700 25697:        final number of relocations from cache: 1205
1701
1702 $ ldd `which evolution` | wc -l
1703      68
1704 $ # Unprelinked system
1705 $ LD_DEBUG=statistics evolution 2>&1 | sed 's/^ *//'
1706 19042:
1707 19042:    runtime linker statistics:
1708 19042:        total startup time in dynamic loader: 54382122 clock cycles
1709 19042:            time needed for relocation: 43403190 clock cycles (79.8%)
1710 19042:            number of relocations: 3452
1711 19042:            number of relocations from cache: 2885
1712 19042:            number of relative relocations: 34957
1713 19042:            time needed to load objects: 10450142 clock cycles (19.2%)
1714
1715 (evolution:19042): Gtk-WARNING **: cannot open display:
1716 19042:
1717 19042:    runtime linker statistics:
1718 19042:        final number of relocations: 4075
1719 19042:        final number of relocations from cache: 2885
1720 $ # Prelinked system
1721 $ LD_DEBUG=statistics evolution 2>&1 | sed 's/^ *//'
1722 25723:
1723 25723:    runtime linker statistics:
1724 25723:        total startup time in dynamic loader: 9176140 clock cycles
1725 25723:            time needed for relocation: 203783 clock cycles (2.2%)
1726 25723:            number of relocations: 0
1727 25723:            number of relocations from cache: 525
1728 25723:            number of relative relocations: 0
1729 25723:            time needed to load objects: 8405157 clock cycles (91.5%)
1730
1731 (evolution:25723): Gtk-WARNING **: cannot open display:
1732 25723:
```

```
1733 25723: runtime linker statistics:  
1734 25723:           final number of relocations: 31  
1735 25723:           final number of relocations from cache: 525  
1736  
1737 $ ldd `which konqueror` | wc -l  
1738     37  
1739 $ # Unprelinked system  
1740 $ LD_DEBUG=statistics konqueror 2>&1 | sed 's/^ *//'  
1741 18979:  
1742 18979: runtime linker statistics:  
1743 18979:   total startup time in dynamic loader: 131985703 clock cycles  
1744 18979:   time needed for relocation: 127341077 clock cycles (96.4%)  
1745 18979:   number of relocations: 25473  
1746 18979:   number of relocations from cache: 53594  
1747 18979:   number of relative relocations: 31171  
1748 18979:   time needed to load objects: 4318803 clock cycles (3.2%)  
1749 konqueror: cannot connect to X server  
1750 18979:  
1751 18979: runtime linker statistics:  
1752 18979:   final number of relocations: 25759  
1753 18979:   final number of relocations from cache: 53594  
1754 $ # Prelinked system  
1755 $ LD_DEBUG=statistics konqueror 2>&1 | sed 's/^ *//'  
1756 25733:  
1757 25733: runtime linker statistics:  
1758 25733:   total startup time in dynamic loader: 5533696 clock cycles  
1759 25733:   time needed for relocation: 1941489 clock cycles (35.0%)  
1760 25733:   number of relocations: 0  
1761 25733:   number of relocations from cache: 2066  
1762 25733:   number of relative relocations: 0  
1763 25733:   time needed to load objects: 3217736 clock cycles (58.1%)  
1764 konqueror: cannot connect to X server  
1765 25733:  
1766 25733: runtime linker statistics:  
1767 25733:   final number of relocations: 0  
1768 25733:   final number of relocations from cache: 2066  
1769  
1770 $ ldd `which kword` | wc -l  
1771     40  
1772 $ # Unprelinked system  
1773 $ LD_DEBUG=statistics kword 2>&1 | sed 's/^ *//'  
1774 19065:  
1775 19065: runtime linker statistics:  
1776 19065:   total startup time in dynamic loader: 153684591 clock cycles  
1777 19065:   time needed for relocation: 148255294 clock cycles (96.4%)  
1778 19065:   number of relocations: 26231  
1779 19065:   number of relocations from cache: 55833  
1780 19065:   number of relative relocations: 30660  
1781 19065:   time needed to load objects: 5068746 clock cycles (3.2%)  
1782 kword: cannot connect to X server  
1783 19065:  
1784 19065: runtime linker statistics:  
1785 19065:   final number of relocations: 26528  
1786 19065:   final number of relocations from cache: 55833  
1787 $ # Prelinked system  
1788 $ LD_DEBUG=statistics kword 2>&1 | sed 's/^ *//'  
1789 25749:  
1790 25749: runtime linker statistics:  
1791 25749:   total startup time in dynamic loader: 6516635 clock cycles  
1792 25749:   time needed for relocation: 2106856 clock cycles (32.3%)  
1793 25749:   number of relocations: 0  
1794 25749:   number of relocations from cache: 2130  
1795 25749:   number of relative relocations: 0  
1796 25749:   time needed to load objects: 4008585 clock cycles (61.5%)  
1797 kword: cannot connect to X server
```

```

1798 25749:
1799 25749: runtime linker statistics:
1800 25749:           final number of relocations: 0
1801 25749:           final number of relocations from cache: 2130

```

Listing 24: Dynamic linker statistics for unprelinked and prelinked GUI programs

1802 In the case of above mentioned `Gtk+` applications, the original startup time spent in the dynamic linker decreased into
 1803 11% to 17% of the original times, with `KDE` applications it decreased even into around 4.2% of original times.

1804 The startup time reported by the dynamic linker is only part of the total startup time of a GUI program. Unfortunately it
 1805 cannot be measured very accurately without patching each application separately, so that it would print current process
 1806 CPU time at the point when all windows are painted and the process starts waiting for user input. The following table
 1807 contains values reported by `time(1)` command on each of the 4 GUI programs running under X, both on unprelinked
 1808 and fully prelinked system. As soon as each program painted its windows, it was killed by application's quit hot key
 1809 ²¹. Especially the real time values depend also on the speed of human reactions, so each measurement was repeated
 1810 10 times. All timings were done with hot caches, after running the applications two times before measurement.

| Type | Values (in seconds) | | | | | | | | | | Average | Std.Dev. |
|-----------------------|---------------------|-------|-------|-------|-------|-------|-------|-------|-------|-------|---------|----------|
| unprelinked epiphany | | | | | | | | | | | | |
| real | 3.053 | 2.84 | 2.996 | 2.901 | 3.019 | 2.929 | 2.883 | 2.975 | 2.922 | 3.026 | 2.954 | 0.0698 |
| user | 2.33 | 2.31 | 2.28 | 2.32 | 2.44 | 2.37 | 2.29 | 2.35 | 2.34 | 2.41 | 2.344 | 0.0508 |
| sys | 0.2 | 0.23 | 0.23 | 0.19 | 0.19 | 0.12 | 0.25 | 0.16 | 0.14 | 0.14 | 0.185 | 0.0440 |
| prelinked epiphany | | | | | | | | | | | | |
| real | 2.773 | 2.743 | 2.833 | 2.753 | 2.753 | 2.644 | 2.717 | 2.897 | 2.68 | 2.761 | 2.755 | 0.0716 |
| user | 2.18 | 2.17 | 2.17 | 2.12 | 2.23 | 2.26 | 2.13 | 2.17 | 2.15 | 2.15 | 2.173 | 0.0430 |
| sys | 0.13 | 0.15 | 0.18 | 0.15 | 0.11 | 0.04 | 0.18 | 0.14 | 0.1 | 0.15 | 0.133 | 0.0416 |
| unprelinked evolution | | | | | | | | | | | | |
| real | 2.106 | 1.886 | 1.828 | 2.12 | 1.867 | 1.871 | 2.242 | 1.871 | 1.862 | 2.241 | 1.989 | 0.1679 |
| user | 1.12 | 1.09 | 1.15 | 1.19 | 1.17 | 1.23 | 1.15 | 1.11 | 1.17 | 1.14 | 1.152 | 0.0408 |
| sys | 0.1 | 0.11 | 0.13 | 0.07 | 0.1 | 0.05 | 0.11 | 0.11 | 0.09 | 0.08 | 0.095 | 0.0232 |
| prelinked evolution | | | | | | | | | | | | |
| real | 1.684 | 1.621 | 1.686 | 1.72 | 1.694 | 1.691 | 1.631 | 1.697 | 1.668 | 1.535 | 1.663 | 0.0541 |
| user | 0.92 | 0.87 | 0.92 | 0.95 | 0.79 | 0.86 | 0.94 | 0.87 | 0.89 | 0.86 | 0.887 | 0.0476 |
| sys | 0.06 | 0.1 | 0.06 | 0.05 | 0.11 | 0.08 | 0.07 | 0.1 | 0.12 | 0.07 | 0.082 | 0.0239 |
| unprelinked kword | | | | | | | | | | | | |
| real | 2.111 | 1.414 | 1.36 | 1.356 | 1.259 | 1.383 | 1.28 | 1.321 | 1.252 | 1.407 | 1.414 | 0.2517 |
| user | 1.04 | 0.9 | 0.93 | 0.88 | 0.89 | 0.89 | 0.87 | 0.89 | 0.9 | 0.8 | 0.899 | 0.0597 |
| sys | 0.07 | 0.04 | 0.06 | 0.05 | 0.06 | 0.1 | 0.09 | 0.08 | 0.08 | 0.12 | 0.075 | 0.0242 |
| prelinked kword | | | | | | | | | | | | |
| real | 1.59 | 1.052 | 0.972 | 1.064 | 1.106 | 1.087 | 1.066 | 1.087 | 1.065 | 1.005 | 1.109 | 0.1735 |
| user | 0.61 | 0.53 | 0.58 | 0.6 | 0.6 | 0.58 | 0.59 | 0.61 | 0.57 | 0.6 | 0.587 | 0.0241 |
| sys | 0.08 | 0.08 | 0.06 | 0.06 | 0.03 | 0.07 | 0.06 | 0.03 | 0.06 | 0.04 | 0.057 | 0.0183 |
| unprelinked konqueror | | | | | | | | | | | | |
| real | 1.306 | 1.386 | 1.27 | 1.243 | 1.227 | 1.286 | 1.262 | 1.322 | 1.345 | 1.332 | 1.298 | 0.0495 |
| user | 0.88 | 0.86 | 0.88 | 0.9 | 0.87 | 0.83 | 0.83 | 0.86 | 0.86 | 0.89 | 0.866 | 0.0232 |
| sys | 0.07 | 0.11 | 0.12 | 0.1 | 0.12 | 0.08 | 0.13 | 0.12 | 0.09 | 0.08 | 0.102 | 0.0210 |
| prelinked konqueror | | | | | | | | | | | | |
| real | 1.056 | 0.962 | 0.961 | 0.906 | 0.927 | 0.923 | 0.933 | 0.958 | 0.955 | 1.142 | 0.972 | 0.0722 |
| user | 0.56 | 0.6 | 0.56 | 0.52 | 0.57 | 0.58 | 0.5 | 0.57 | 0.61 | 0.55 | 0.562 | 0.0334 |
| sys | 0.1 | 0.13 | 0.08 | 0.15 | 0.07 | 0.09 | 0.09 | 0.09 | 0.1 | 0.08 | 0.098 | 0.0244 |

Table 1: GUI program start up times without and with prelinking

1811

²¹Ctrl+W for Epiphany, Ctrl+Q for Evolution and Konqueror and Enter in Kword's document type choice dialog.

1812 OpenOffice.org is probably the largest program these days in Linux, mostly written in C++. In OpenOffice.org
1813 1.1, the main executable, `soffice.bin`, links directly against 34 shared libraries, but typically during startup it loads
1814 using `dlopen` many others. As has been mentioned earlier, `prelink` cannot speed up loading shared libraries using
1815 `dlopen`, since it cannot predict in which order and what shared libraries will be loaded (and thus cannot compute
1816 conflict fixups). The `soffice.bin` is typically started through a wrapper script and depending on what arguments
1817 are passed to it, different OpenOffice.org application is started. With no options, it starts just empty window with
1818 menu from which the applications can be started, with say `private:factory/swriter` argument it starts a word
1819 processor, with `private:factory/scalc` it starts a spreadsheet etc. When `soffice.bin` is already running, if you
1820 start another copy of it, it just instructs the already running copy to pop up a new window and exits.

1821 In an experiment, `soffice.bin` has been invoked 7 times against running X server, with no arguments, `private:factory/swriter`,
1822 `private:factory/scalc`, `private:factory/sdraw`, `private:factory/simpress`, `private:factory/smath`
1823 arguments (in all these cases nothing was pressed at all) and last with the `private:factory/swriter` argument
1824 where the menu item New Presentation was selected and the word processor window closed. In all these cases,
1825 `/proc/'pidof soffice.bin'/maps` file was captured and the application then killed. This file contains among
1826 other things list of all shared libraries mmaped by the process at the point where it started waiting for user input
1827 after loading up. These lists were then summarized, to get number of the runs in which particular shared library was
1828 loaded up out of the total 7 runs. There were 38 shared libraries shipped as part of OpenOffice.org package which
1829 have been loaded in all 7 times, another 3 shared libraries included in OpenOffice.org (and also one shared library
1830 shipped in another package, `libdb_cxx-4.1.so`) which were loaded 6 times.²² There was one shared library loaded
1831 in 5 runs, but was locale specific and thus not worth considering. Inspecting OpenOffice.org source, these shared
1832 libraries are never unloaded with `dlclose`, so `soffice.bin` can be made much more `prelink` friendly and thus save
1833 substantial amount of startup time by linking against all those 76 shared libraries instead of just 34 shared libraries it is
1834 linked against. In the timings below, `soffice1.bin` is the original `soffice.bin` as created by the OpenOffice.org
1835 makefiles and `soffice3.bin` is the same executable linked dynamically against additional 42 shared libraries. The
1836 ordering of those 42 shared libraries matters for the number of conflict fixups, unfortunately with large C++ shared
1837 libraries there is no obvious rule for ordering them as sometimes it is more useful when a shared library precedes its
1838 dependency and sometimes vice versa, so a few different orderings were tried in several steps and always the one with
1839 smallest number of conflict fixups was chosen. Still, the number of conflict fixups is quite high and big part of the
1840 fixups are storing addresses of PLT slots in the executable into various places in shared libraries²³ `soffice2.bin` is
1841 another experiment, where the executable itself is empty source file, all objects which were originally in `soffice.bin`
1842 executable with the exception of start files were recompiled as position independent code and linked into a new shared
1843 library. This reduced number of conflicts a lot and speeded up start up times against `soffice3.bin` when caches are
1844 hot. It is a little bit slower than `soffice3.bin` when running with cold caches (e.g. for the first time after bootup), as
1845 there is one more shared library to load etc.

1846 In the timings below, numbers for `soffice1.bin` and `soffice2.bin` resp. `soffice3.bin` cannot be easily com-
1847 pared, as `soffice1.bin` loads less than half of the needed shared libraries which the remaining two executables load
1848 and the time to load those shared libraries doesn't show up there. Still, when it is prelinked it takes just slightly more
1849 than two times longer to load `soffice2.bin` than `soffice1.bin` and the times are still less than 7% of how long it
1850 takes to load just the initial 34 shared libraries when not prelinking.

```
1851 $ S='s/^ *//'
1852 $ ldd /usr/lib/openoffice/program/soffice1.bin | wc -l
1853      34
1854 $ # Unprelinked system
1855 $ LD_DEBUG=statistics /usr/lib/openoffice/program/soffice1.bin 2>&1 | sed "$S"
1856 19095:
1857 19095: runtime linker statistics:
1858 19095:   total startup time in dynamic loader: 159833582 clock cycles
1859 19095:           time needed for relocation: 155464174 clock cycles (97.2%)
1860 19095:           number of relocations: 31136
1861 19095:           number of relocations from cache: 31702
1862 19095:           number of relative relocations: 18284
1863 19095:           time needed to load objects: 3919645 clock cycles (2.4%)
```

²²In all runs but when ran without arguments. But when the application is started without any arguments, it cannot do any useful work, so one loads one of the applications afterward anyway.

²³This might get better when the linker is modified to handle calls without ever taking address of the function in executables specially, but only testing it will actually show it up.

```
1864 /usr/lib/openoffice/program/soffice1.bin X11 error: Can't open display:  
1865 Set DISPLAY environment variable, use -display option  
1866 or check permissions of your X-Server  
1867 (See "man X" resp. "man xhost" for details)  
1868 19095:  
1869 19095: runtime linker statistics:  
1870 19095:           final number of relocations: 31715  
1871 19095: final number of relocations from cache: 31702  
1872 $ # Prelinked system  
1873 $ LD_DEBUG=statistics /usr/lib/openoffice/program/soffice1.bin 2>&1 | sed "$S"  
1874 25759:  
1875 25759: runtime linker statistics:  
1876 25759: total startup time in dynamic loader: 4252397 clock cycles  
1877 25759:           time needed for relocation: 1189840 clock cycles (27.9%)  
1878 25759:           number of relocations: 0  
1879 25759: number of relocations from cache: 2142  
1880 25759:           number of relative relocations: 0  
1881 25759:           time needed to load objects: 2604486 clock cycles (61.2%)  
1882 /usr/lib/openoffice/program/soffice1.bin X11 error: Can't open display:  
1883 Set DISPLAY environment variable, use -display option  
1884 or check permissions of your X-Server  
1885 (See "man X" resp. "man xhost" for details)  
1886 25759:  
1887 25759: runtime linker statistics:  
1888 25759:           final number of relocations: 24  
1889 25759: final number of relocations from cache: 2142  
1890 $ ldd /usr/lib/openoffice/program/soffice2.bin | wc -l  
1891 77  
1892 $ # Unprelinked system  
1893 $ LD_DEBUG=statistics /usr/lib/openoffice/program/soffice2.bin 2>&1 | sed "$S"  
1894 19115:  
1895 19115: runtime linker statistics:  
1896 19115: total startup time in dynamic loader: 947793670 clock cycles  
1897 19115:           time needed for relocation: 936895741 clock cycles (98.8%)  
1898 19115:           number of relocations: 69164  
1899 19115: number of relocations from cache: 94502  
1900 19115:           number of relative relocations: 59374  
1901 19115:           time needed to load objects: 10046486 clock cycles (1.0%)  
1902 /usr/lib/openoffice/program/soffice2.bin X11 error: Can't open display:  
1903 Set DISPLAY environment variable, use -display option  
1904 or check permissions of your X-Server  
1905 (See "man X" resp. "man xhost" for details)  
1906 19115:  
1907 19115: runtime linker statistics:  
1908 19115:           final number of relocations: 69966  
1909 19115: final number of relocations from cache: 94502  
1910 $ # Prelinked system  
1911 $ LD_DEBUG=statistics /usr/lib/openoffice/program/soffice2.bin 2>&1 | sed "$S"  
1912 25777:  
1913 25777: runtime linker statistics:  
1914 25777: total startup time in dynamic loader: 10952099 clock cycles  
1915 25777:           time needed for relocation: 3254518 clock cycles (29.7%)  
1916 25777:           number of relocations: 0  
1917 25777: number of relocations from cache: 5309  
1918 25777:           number of relative relocations: 0  
1919 25777:           time needed to load objects: 6805013 clock cycles (62.1%)  
1920 /usr/lib/openoffice/program/soffice2.bin X11 error: Can't open display:  
1921 Set DISPLAY environment variable, use -display option  
1922 or check permissions of your X-Server  
1923 (See "man X" resp. "man xhost" for details)  
1924 25777:  
1925 25777: runtime linker statistics:  
1926 25777:           final number of relocations: 24  
1927 25777: final number of relocations from cache: 5309
```

```

1928 $ ldd /usr/lib/openoffice/program/soffice3.bin | wc -l
1929      76
1930 $ # Unprelinked system
1931 $ LD_DEBUG=statistics /usr/lib/openoffice/program/soffice3.bin 2>&1 | sed "$S"
1932 19131:
1933 19131: runtime linker statistics:
1934 19131:   total startup time in dynamic loader: 852275754 clock cycles
1935 19131:           time needed for relocation: 840996859 clock cycles (98.6%)
1936 19131:           number of relocations: 68362
1937 19131:           number of relocations from cache: 89213
1938 19131:           number of relative relocations: 55831
1939 19131:           time needed to load objects: 10170207 clock cycles (1.1%)
1940 /usr/lib/openoffice/program/soffice3.bin X11 error: Can't open display:
1941 Set DISPLAY environment variable, use -display option
1942 or check permissions of your X-Server
1943 (See "man X" resp. "man xhost" for details)
1944 19131:
1945 19131: runtime linker statistics:
1946 19131:           final number of relocations: 69177
1947 19131:           final number of relocations from cache: 89213
1948 $ # Prelinked system
1949 $ LD_DEBUG=statistics /usr/lib/openoffice/program/soffice3.bin 2>&1 | sed "$S"
1950 25847:
1951 25847: runtime linker statistics:
1952 25847:   total startup time in dynamic loader: 12277407 clock cycles
1953 25847:           time needed for relocation: 4232915 clock cycles (34.4%)
1954 25847:           number of relocations: 0
1955 25847:           number of relocations from cache: 8961
1956 25847:           number of relative relocations: 0
1957 25847:           time needed to load objects: 6925023 clock cycles (56.4%)
1958 /usr/lib/openoffice/program/soffice3.bin X11 error: Can't open display:
1959 Set DISPLAY environment variable, use -display option
1960 or check permissions of your X-Server
1961 (See "man X" resp. "man xhost" for details)
1962 25847:
1963 25847: runtime linker statistics:
1964 25847:           final number of relocations: 24
1965 25847:           final number of relocations from cache: 8961

```

Listing 25: Dynamic linker statistics for unprelinked and prelinked OpenOffice.org

1966 Below are measurement using `time(1)` for each of the `soffice.bin` variants, prelinked and unprelinked. OpenOffice.org
1967 was killed immediately after painting Writer's window using `Ctrl+Q`.

| Type | Values (in seconds) | | | | | | | | | | Average | Std.Dev. |
|--|---------------------|-------|-------|-------|-------|-------|-------|-------|-------|-------|---------|----------|
| unprelinked soffice1.bin private:factory/swriter | | | | | | | | | | | | |
| real | 5.569 | 5.149 | 5.547 | 5.559 | 5.549 | 5.139 | 5.55 | 5.559 | 5.598 | 5.559 | 5.478 | 0.1765 |
| user | 4.65 | 4.57 | 4.62 | 4.64 | 4.57 | 4.55 | 4.65 | 4.49 | 4.52 | 4.46 | 4.572 | 0.0680 |
| sys | 0.29 | 0.24 | 0.19 | 0.21 | 0.21 | 0.21 | 0.25 | 0.25 | 0.27 | 0.26 | 0.238 | 0.0319 |
| prelinked soffice1.bin private:factory/swriter | | | | | | | | | | | | |
| real | 4.946 | 4.899 | 5.291 | 4.879 | 4.879 | 4.898 | 5.299 | 4.901 | 4.887 | 4.901 | 4.978 | 0.1681 |
| user | 4.23 | 4.27 | 4.18 | 4.24 | 4.17 | 4.22 | 4.15 | 4.25 | 4.26 | 4.31 | 4.228 | 0.0494 |
| sys | 0.22 | 0.22 | 0.24 | 0.26 | 0.3 | 0.26 | 0.29 | 0.17 | 0.21 | 0.23 | 0.24 | 0.0389 |
| unprelinked soffice2.bin private:factory/swriter | | | | | | | | | | | | |
| real | 5.575 | 5.166 | 5.592 | 5.149 | 5.571 | 5.559 | 5.159 | 5.157 | 5.569 | 5.149 | 5.365 | 0.2201 |
| user | 4.59 | 4.5 | 4.57 | 4.37 | 4.47 | 4.57 | 4.56 | 4.41 | 4.63 | 4.5 | 4.517 | 0.0826 |
| sys | 0.24 | 0.24 | 0.21 | 0.34 | 0.27 | 0.19 | 0.19 | 0.27 | 0.19 | 0.29 | 0.243 | 0.0501 |
| prelinked soffice2.bin private:factory/swriter | | | | | | | | | | | | |
| real | 3.69 | 3.66 | 3.658 | 3.661 | 3.639 | 3.638 | 3.649 | 3.659 | 3.65 | 3.659 | 3.656 | 0.0146 |
| user | 2.93 | 2.88 | 2.88 | 2.9 | 2.84 | 2.63 | 2.89 | 2.85 | 2.77 | 2.83 | 2.84 | 0.0860 |
| sys | 0.22 | 0.18 | 0.23 | 0.2 | 0.18 | 0.29 | 0.22 | 0.23 | 0.24 | 0.22 | 0.221 | 0.0318 |

| Type | Values (in seconds) | | | | | | | | | | Average | Std.Dev. |
|------|--|-------|-------|-------|-------|-------|-------|-------|-------|-------|---------|----------|
| | unprelinked soffice3.bin private:factory/swriter | | | | | | | | | | | |
| real | 5.031 | 5.02 | 5.009 | 5.028 | 5.019 | 5.019 | 5.019 | 5.052 | 5.426 | 5.029 | 5.065 | 0.1273 |
| user | 4.31 | 4.35 | 4.34 | 4.3 | 4.38 | 4.29 | 4.45 | 4.37 | 4.38 | 4.44 | 4.361 | 0.0547 |
| sys | 0.27 | 0.25 | 0.26 | 0.27 | 0.27 | 0.31 | 0.18 | 0.17 | 0.16 | 0.15 | 0.229 | 0.0576 |
| | prelinked soffice3.bin private:factory/swriter | | | | | | | | | | | |
| real | 3.705 | 3.669 | 3.659 | 3.669 | 3.66 | 3.659 | 3.659 | 3.661 | 3.668 | 3.649 | 3.666 | 0.0151 |
| user | 2.86 | 2.88 | 2.85 | 2.84 | 2.83 | 2.86 | 2.84 | 2.91 | 2.86 | 2.8 | 2.853 | 0.0295 |
| sys | 0.26 | 0.19 | 0.27 | 0.25 | 0.24 | 0.23 | 0.28 | 0.21 | 0.21 | 0.27 | 0.241 | 0.0303 |

Table 2: OpenOffice.org start up times without and with prelinking

1968

15 Similar tools on other ELF using Operating Systems

1969 Something similar to `prelink` is available on other ELF platforms. On Irix there is `QUICKSTART` and on Solaris `crle`.

1970 SGI `QUICKSTART` is much closer to `prelink` from these two. The `rqs` program relocates libraries to (if possible)
 1971 unique virtual address space slot. The base address is either specified on the command line with the `-l` option, or `rqs`
 1972 uses a `so.locations` registry with `-c` or `-u` options and finds a not yet occupied slot. This is similar to how `prelink`
 1973 lays out libraries without the `-m` option.

1974 `QUICKSTART` uses the same data structure for library lists (`ElfNN_Lib`) as `prelink`, but uses more fields in it
 1975 (`prelink` doesn't use `l_version` and `l_flags` fields at the moment) and uses different dynamic tags and section
 1976 type for it. Another difference is that `QUICKSTART` makes all liblist section `SHF_ALLOC`, whether in shared libraries or
 1977 executables. `prelink` only needs liblist section in the executable be allocated, liblist sections in shared libraries are
 1978 not allocated and used at `prelink` time only.

1979 The biggest difference between `QUICKSTART` and `prelink` is in how conflicts are encoded. SGI stores them in a
 1980 very compact format, as array of `.dynsym` section indexes for symbols which are conflicting. There is no information
 1981 publicly available what exactly SGI dynamic linker does when it is resolving the conflicts, so this is just a guess. Given
 1982 that the conflicts can be stored in a shared library or executable different to the shared library with the relocations
 1983 against the conflicting symbol and different to the shared library which the symbol was originally resolved to, there
 1984 doesn't seem to be an obvious way how to handle the conflicts very cheaply. The dynamic linker probably collects
 1985 list of all conflicting symbol names, for each such symbol computes ELF hash and walks hash buckets for this hash
 1986 of all shared libraries, looking for the symbol. Every time it finds the symbol, all relocations against it need to be
 1987 redone. Unlike this, `prelink` stores conflicts as an array of `ElfNN_RelA` structures, with one entry for each shared
 1988 relocation against conflicting symbol in some shared library. This guarantees that there are no symbol lookups during
 1989 program startup (provided that shared libraries have not been changed after prelinking), while with `QUICKSTART` will
 1990 do some symbol lookups if there are any conflicts. `QUICKSTART` puts conflict sections into the executable and every
 1991 shared library where `rqs` determines conflicts while `prelink` stores them in the executable only (but the array is
 1992 typically much bigger). Disk space requirements for prelinked executables are certainly bigger than for requickstarted
 1993 executables, but which one has bigger runtime memory requirements is unclear. If prelinking can be used, all `.rela*`
 1994 and `.rel*` sections in the executable and all shared libraries are skipped, so they will not need to be paged in during
 1995 whole program's life (with the exception of first and last pages in the relocation sections which can be paged in because
 1996 of other sections on the same page), but whole `.gnu.conflict` section needs to be paged in (read-only) and processed.
 1997 With `QUICKSTART`, probably all (much smaller) conflict sections need to be paged in and also likely for each conflict
 1998 whole relocation sections of each library which needs the conflict to be applied against.

1999 In `QUICKSTART` documentation, SGI says that conflicts are very costly and that developers should avoid them. Un-
 2000 fortunately, this is sometimes quite hard, especially with C++ shared libraries. It is unclear whether `rqs` does any
 2001 optimizations to trim down the number of conflicts.

2002 Sun took completely different approach. The dynamic linker provides a `dldump (const char *ipath, const`
 2003 `char *opath, int flags);` function. `ipath` is supposed to be a path to an ELF object loaded already in the current
 2004 process. This function creates a new ELF object at `opath`, which is like the `ipath` object, but relocated to the base address
 2005 which it has actually been mapped at in the current process and with some relocations (specified in `flags` bitmask)

2006 applied as they have been resolved in the current process. Relocations, which have been applied, are overwritten in
 2007 the relocation sections with R_*_NONE relocations. The crle executable, in addition to other functions not related to
 2008 startup times, with some specific options uses the dldump function to dump all shared libraries a particular executable
 2009 uses (and the executable itself) into a new directory, with selected relocation classes being already applied. The main
 2010 disadvantage of this approach is that such alternate shared libraries are at least for most relocation classes not shareable
 2011 across different programs at all (and for those where they could be shareable a little bit there will be many relocations
 2012 left for the dynamic linker, so the speed gains will be small). Another disadvantage is that all relocation sections need
 2013 to be paged into the memory, just to find out that most of the relocations are R_*_NONE.

16 ELF extensions for prelink

2014 Prelink needs a few ELF extensions for its data structures in ELF objects. For list of dependencies at the time of
 2015 prelinking, a new section type SHT_GNU_LIBLIST is defined:

```
2016 #define SHT_GNU_LIBLIST 0x6fffffff /* Prelink library list */
2017
2018 typedef struct
2019 {
2020   Elf32_Word l_name;           /* Name (string table index) */
2021   Elf32_Word l_time_stamp;    /* Timestamp */
2022   Elf32_Word l_checksum;      /* Checksum */
2023   Elf32_Word l_version;      /* Unused, should be zero */
2024   Elf32_Word l_flags;         /* Unused, should be zero */
2025 } Elf32_Lib;
2026
2027 typedef struct
2028 {
2029   Elf64_Word l_name;           /* Name (string table index) */
2030   Elf64_Word l_time_stamp;    /* Timestamp */
2031   Elf64_Word l_checksum;      /* Checksum */
2032   Elf64_Word l_version;      /* Unused, should be zero */
2033   Elf64_Word l_flags;         /* Unused, should be zero */
2034 } Elf64_Lib;
```

Listing 26: New structures and section type constants used by prelink

2035 Introduces a few new special sections:

| Name | Type | Attributes |
|----------------------------|-----------------|------------|
| <i>In shared libraries</i> | | |
| .gnu.liblist | SHT_GNU_LIBLIST | 0 |
| .gnu.libstr | SHT_STRTAB | 0 |
| .gnu.prelink_undo | SHT_PROGBITS | 0 |
| <i>In executables</i> | | |
| .gnu.liblist | SHT_GNU_LIBLIST | SHF_ALLOC |
| .gnu.conflict | SHT_RELATIVE | SHF_ALLOC |
| .gnu.prelink_undo | SHT_PROGBITS | 0 |

Table 3: Special sections introduced by prelink

2036

2037 .gnu.liblist This section contains one ElfNN_Lib structure for each shared library which the object has been pre-

2038 linked against, in the order in which they appear in symbol search scope. Section's `sh_link` value should contain
2039 section index of `.gnu.libstr` for shared libraries and section index of `.dynsym` for executables. `l_name` field
2040 contains the dependent library's name as index into the section pointed by `sh_link` field. `l_time_stamp` resp.
2041 `l_checksum` should contain copies of `DT_GNU_PRELINKED` resp. `DT_CHECKSUM` values of the dependent library.

2042 **.gnu.conflict** This section contains one `ElfNN_Rela` structure for each needed prelink conflict fixup. `r_offset`
2043 field contains the absolute address at which the fixup needs to be applied, `r_addend` the value that needs to be
2044 stored at that location. `ELFNN_R_SYM` of `r_info` field should be zero, `ELFNN_R_TYPE` of `r_info` field should be
2045 architecture specific relocation type which should be handled the same as for `.rela.*` sections on the archi-
2046 tecture. For EM_ALPHA machine, all types with `R_ALPHA_JMP_SLOT` in lowest 8 bits of `ELF64_R_TYPE` should be
2047 handled as `R_ALPHA_JMP_SLOT` relocation, the upper 24 bits contains index in original `.rela.plt` section of the
2048 `R_ALPHA_JMP_SLOT` relocation the fixup was created for.

2049 **.gnu.libstr** This section contains strings for `.gnu.liblist` section in shared libraries where `.gnu.liblist`
2050 section is not allocated.

2051 **.gnu.prelink_undo** This section contains prelink private data used for `prelink --undo` operation. This data in-
2052 cludes the original `ElfNN_Ehdr` of the object before prelinking and all its original `ElfNN_Phdr` and `ElfNN_Shdr`
2053 headers.

2054 Prelink also defines 6 new dynamic tags:

```
2055 #define DT_GNU_PRELINKED 0x6fffffdf5 /* Prelinking timestamp */  
2056 #define DT_GNU_CONFLICTSZ 0x6fffffdf6 /* Size of conflict section */  
2057 #define DT_GNU_LIBLISTSZ 0x6fffffdf7 /* Size of library list */  
2058 #define DT_CHECKSUM 0x6fffffdf8 /* Library checksum */  
2059  
2060 #define DT_GNU_CONFLICT 0x6fffffef8 /* Start of conflict section */  
2061 #define DT_GNU_LIBLIST 0x6fffffef9 /* Library list */
```

Listing 27: Prelink dynamic tags

2062 `DT_GNU_PRELINKED` and `DT_CHECKSUM` dynamic tags must be present in prelinked shared libraries. The corresponding
2063 `d_un.d_val` fields should contain time when the library has been prelinked (in seconds since January, 1st, 1970, 00:00
2064 UTC) resp. CRC32 checksum of all sections with one of `SHF_ALLOC`, `SHF_WRITE` or `SHF_EXECINSTR` bit set whose
2065 type is not `SHT_NOBITS`, in the order they appear in the shared library's section header table, with `DT_GNU_PRELINKED`
2066 and `DT_CHECKSUM` `d_un.v_val` values set to 0 for the time of checksum computation.

2067 The `DT_GNU_LIBLIST` and `DT_GNU_LIBLISTSZ` dynamic tags must be present in all prelinked executables. The
2068 `d_un.d_ptr` value of the `DT_GNU_LIBLIST` dynamic tag contains the virtual address of the `.gnu.liblist` section
2069 in the executable and `d_un.d_val` of `DT_GNU_LIBLISTSZ` tag contains its size in bytes.

2070 `DT_GNU_CONFLICT` and `DT_GNU_CONFLICTSZ` dynamic tags may be present in prelinked executables. `d_un.d_ptr` of
2071 `DT_GNU_CONFLICT` dynamic tag contains the virtual address of `.gnu.conflict` section in the executable (if present)
2072 and `d_un.d_val` of `DT_GNU_CONFLICTSZ` tag contains its size in bytes.

A Glossary

Nomenclature

2074 ASCII Shield area First 16MB of address space on 32-bit architectures. These addresses have zeros in upper 8 bits,
2075 which on little endian architectures are stored as last byte of the address and on big endian architectures as first
2076 byte of the address. A zero byte terminates string, so it is hard to control the exact arguments of a function if
2077 they are placed on the stack above the address. On big endian machines, it is even hard to control the low 24
2078 bits of the address,

2079 Global Offset Table (GOT) When position independent code needs to build address which requires dynamic relocation,
2080 instead of building it as constant in registers and applying a dynamic relocation against the read-only segment
2081 (which would mean that any pages of the read-only segment where relocations are applied cannot be shared
2082 between processes anymore), it loads the address from an offset table private to each shared library, which
2083 is created by the linker. The table is in writable segment and relocations are applied against it. Position
2084 independent code uses on most architectures a special `PIC` register which points to the start of the Global
2085 Offset Table,

2086 Lazy Binding A way to postpone symbol lookups for calls until a function is called for the first time in particular
2087 shared library. This decreases number of symbol lookups done during startup and symbols which are never
2088 called don't need to be looked up at all. Calls requiring relocations jump into `PLT`, which is initially set up
2089 so that a function in the dynamic linker is called to do symbol lookup. The looked up address is then stored
2090 either into the `PLT` slot directly (if `PLT` is writable) or into GOT entry corresponding to the `PLT` slot and any
2091 subsequent calls already go directly to that address. Lazy binding can be turned off by setting `LD_BIND_NOW=1`
2092 in the environment. Prelinked programs never use lazy binding for the executable or any shared libraries not
2093 loaded using `dlopen`,

2094 Page Memory block of fixed size which virtual memory subsystem deals with as a unit. The size of the page depends
2095 on the addressing hardware of the processor, typically pages are 4K or 8K, in some cases bigger,

2096 PLT Process Linkage Table. Stubs in `ELF` shared libraries and executables which allow lazy relocations of function
2097 calls. They initially point to code which will do the symbol lookup. The result of this symbol lookup is
2098 then stored in the Process Linkage Table and control transferred to the address symbol lookup returned. All
2099 following calls to the `PLT` slot just branch to the already looked up address directly, no further symbol lookup
2100 is needed,

2101 Position Independent Executable A hybrid between classical `ELF` executables and `ELF` shared libraries. It has a form
2102 of a `ET_DYN` object like shared libraries and should contain position independent code, so that the kernel
2103 can load the executable starting at random address to make certain security attacks harder. Unlike shared
2104 libraries it contains `DT_DEBUG` dynamic tag, must have `PT_INTERP` segment with dynamic linker's path, must
2105 have meaningful code at its `e_entry` and can use symbol lookup assumptions normal executables can make,
2106 particularly that no symbol defined in the executable can be overridden by a shared library symbol,

2107 REL Type of relocation structure which includes just offset, relocation type and symbol. Addend is taken from
2108 memory location at offset,

2109 RELA Type of relocation structure which includes offset, relocation type, symbol against which the relocation is and
2110 an integer addend which is added to the symbol. Memory at offset is not supposed to be used by the relocation.
2111 Some architectures got this implemented incorrectly and memory at offset is for some relocation types used by
2112 the relocation, either in addition to addend or addend is not used at all. `RELA` relocations are generally better
2113 for `prelink`, since when `prelink` stores a pre-computed value into the memory location at offset, the addend
2114 value is not lost,

2115 relative relocation Relocation, which doesn't need a symbol lookup, just adds a shared library load offset to certain
2116 memory location (or locations),

2117 RTTI C++ runtime type identification,

2118 Symbol Search Scope The sequence of `ELF` objects in which a symbol is being looked up. When a symbol definition
2119 is found, the searching stops and the found symbol is returned. Each program has a global search scope,
2120 which starts by the executable, is typically followed by the immediate dependencies of the executable and
2121 then their dependencies in breadth search order (where only first occurrence of each shared library is kept).
2122 If `DT_FILTER` or `DT_AUXILIARY` dynamic tags are used the order is slightly different. Each shared library
2123 loaded with `dlopen` has its own symbol search scope which contains that shared library and its dependencies.
2124 `Prelink` operates also with natural symbol search scope of each shared library, which is the global symbol
2125 search scope the shared library would have if it were started as the main program,

B References

2126 [1] System V Application Binary Interface, Edition 4.1.

2127 [2] System V Application Binary Interface, Intel 386 Architecture Processor Supplement.

-
- ₂₁₂₈ [3] *System V Application Binary Interface, AMD64 Architecture Processor Supplement.*
- ₂₁₂₉ [4] *System V Application Binary Interface, Intel Itanium Architecture Processor Supplement*, Intel Corporation, 2001.
- ₂₁₃₀ [5] Steve Zucker, Kari Karhi, *System V Application Binary Interface, PowerPC Architecture Processor Supplement*, SunSoft, IBM, 1995.
- ₂₁₃₂ [6] *System V Application Binary Interface, PowerPC64 Architecture Processor Supplement.*
- ₂₁₃₃ [7] *System V Application Binary Interface, ARM Architecture Processor Supplement.*
- ₂₁₃₄ [8] *SPARC Compliance Definition, Version 2.4.1*, SPARC International, Inc., 1999.
- ₂₁₃₅ [9] Ulrich Drepper, *How To Write Shared Libraries*, Red Hat, Inc., 2003.
- ₂₁₃₆ [10] *Linker And Library Guide*, Sun Microsystems, 2002.
- ₂₁₃₇ [11] John R. Levine, *Linkers and Loaders*, 1999.
- ₂₁₃₈ [12] Ulrich Drepper, *ELF Handling For Thread-Local Storage*, Red Hat, Inc., 2003.
- ₂₁₃₉ [13] Alan Modra, *PowerPC Specific Thread Local Storage ABI*, 2003.
- ₂₁₄₀ [14] Alan Modra, *PowerPC64 Specific Thread Local Storage ABI*, 2003.
- ₂₁₄₁ [15] *DWARF Debugging Information Format Version 2*.
- ₂₁₄₂ [16] *DWARF Debugging Information Format Version 3*, Draft, 2001.
- ₂₁₄₃ [17] The "stabs" debugging information format.

C Revision History

₂₁₄₄ 2003-11-03 First draft.